

ETL design training session – basic

This document serves as a training session for using Talend Open Studio tool (from here on, we refer to it as Talend) for designing an ETL process that answers user analytical needs (expressed in a form of textual requirements).

Content:

Part A: Data vs. Control Flow In the first part of this training, we will briefly explain the difference between data flow (i.e., subjob) and control flow (i.e., job) in Talend and how to create them.

Part B: Create an ETL data flow (transformation) that answers the KPI During the second part, we will use an example textual requirement (KPI 1) to walk you through the process of creating an ETL data flow design using Talend.

Part C: Create control flow (job) to execute an ETL process Lastly, in the third part, we will create a control flow for managing the execution of our ETL process.

Lastly, in Appendix A, we will export the created Talend project into a .zip archive.

Data sources used in the training: All examples are created over the ACME Flying Use Case. For better understanding of internals of the system under the study, the schematic/diagrammatic representations of the domain and the available data sources are available (see slides: **FlyingUseCase.pdf**):

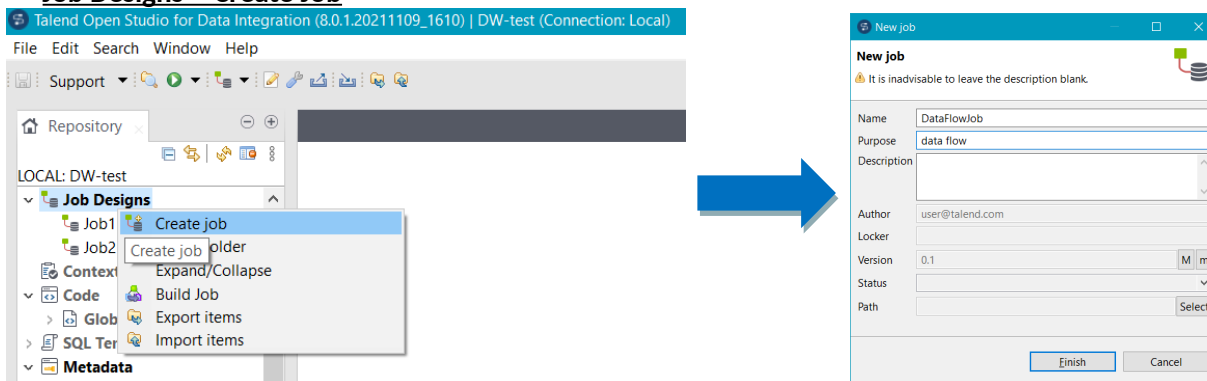
- 1) Diagrammatic/schematic representation of the **ACME subsystems** (AIMS and AMOS)
- 2) AIMS and AMOS DB schema (IE notation)
 - a. With SQL to create **AIMS** and **AMOS tables** (**AIMS.sql** and **AMOS.sql**).
- 3) External sources:
 - a. **aircraft-manufacturerinfo-lookup.csv**
 - b. **maintenance-personnel-airport-lookup.csv**

ETL design: training session - basic

Part A: Data vs. Control Flow

In ETL design, we primarily differentiate between data and control flow.

- **Data flow** is in charge of performing various operations over data themselves in order to prepare them for loading into a DW (or directly for exploiting them by end users). That is, data **extraction** (reading from the data sources), various data **transformation** tasks (data cleaning, integration, format conversions, etc.) and finally **loading** of the data to previously created target data stores of the DW. Data flows are typically executed as **pipelines** of operations (rather than a set of strictly sequential steps started one after another), meaning that they all start executing at the same time but keep idle until they receive data at the input. In Talend, data flows are called **jobs** (or **subjobs if they belong to a control flow job**), and they are created as follows: **Right click on Job Designs-> Create Job**



- **Control flow**, on the other side is responsible of orchestrating the execution of one or more data flows. It does not work directly with data, but rather on managing the execution of data processing (i.e., scheduling, starting, checking for possible errors occurred during the execution, etc.). Unlike data flow, in control flow the order of execution is strictly defined by the sequence of activities, meaning that one activity does not start its execution until all its input activities have finished. This is especially important in the case of dependent data processing, where the results of one data flow are needed before starting the execution of another. In Talend, control flows are also called **jobs**, but they consist of several subjobs each one representing one data flow. Control flows are created within the same environment as data flows, but by using the **tRunJob**¹ component from the **Orchestration** palette. Each of these components references some previously created subjob (data flow) and in combination with other components from the **Orchestration** palette² creates a control flow job.

Part B: Create an ETL data flow (transformation) that answers the KPI

In this exercise, we will simulate the creation of an example ETL process with a single KPI in mind. Notice that while ETLs typically require several complex data flows (managing the loading of both factual and dimensional tables), here we practice on a single data flow that prepares data to answer a single user requirement (i.e., to create a multidimensional cube that can answer such KPI).

KPI 1

“Analyze the total number of **flight hours** (*actualArrival - actualDeparture*) of an aircraft from the point of view of the **aircraft's model** (*aircraft_model*) and **year** (*extract(year from actualDeparture)*) of the flight with **number of passengers** (*passengers*) greater than 50. “

¹ <https://help.talend.com/r/en-US/8.0/orchestration/trunjob>

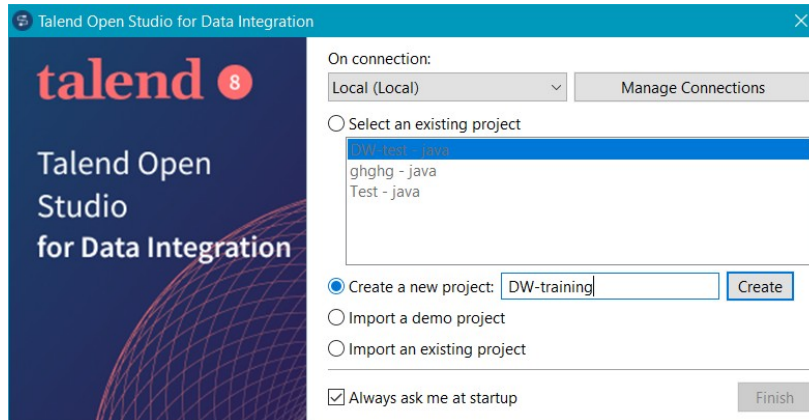
² <https://help.talend.com/r/en-US/8.0/orchestration/orchestration>

Design process:

The ETL design in Talend should be provided using their graphical editor (**Talend Open Studio for Data Integration**).

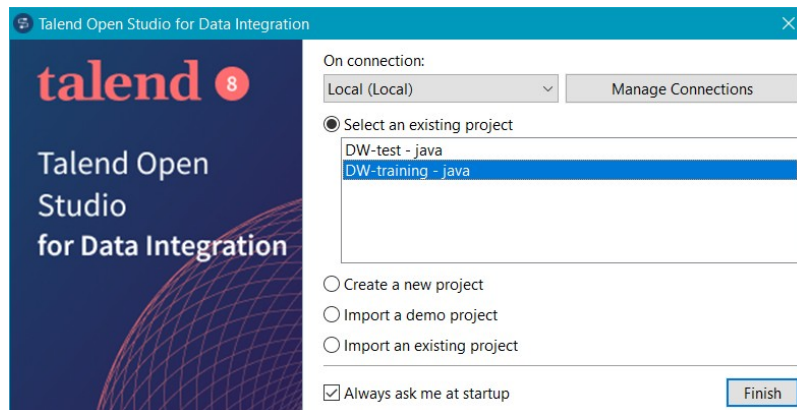
1. Opening Talend's graphical ETL editor and creating new Talend project

After we open Talend, we first see the following screen



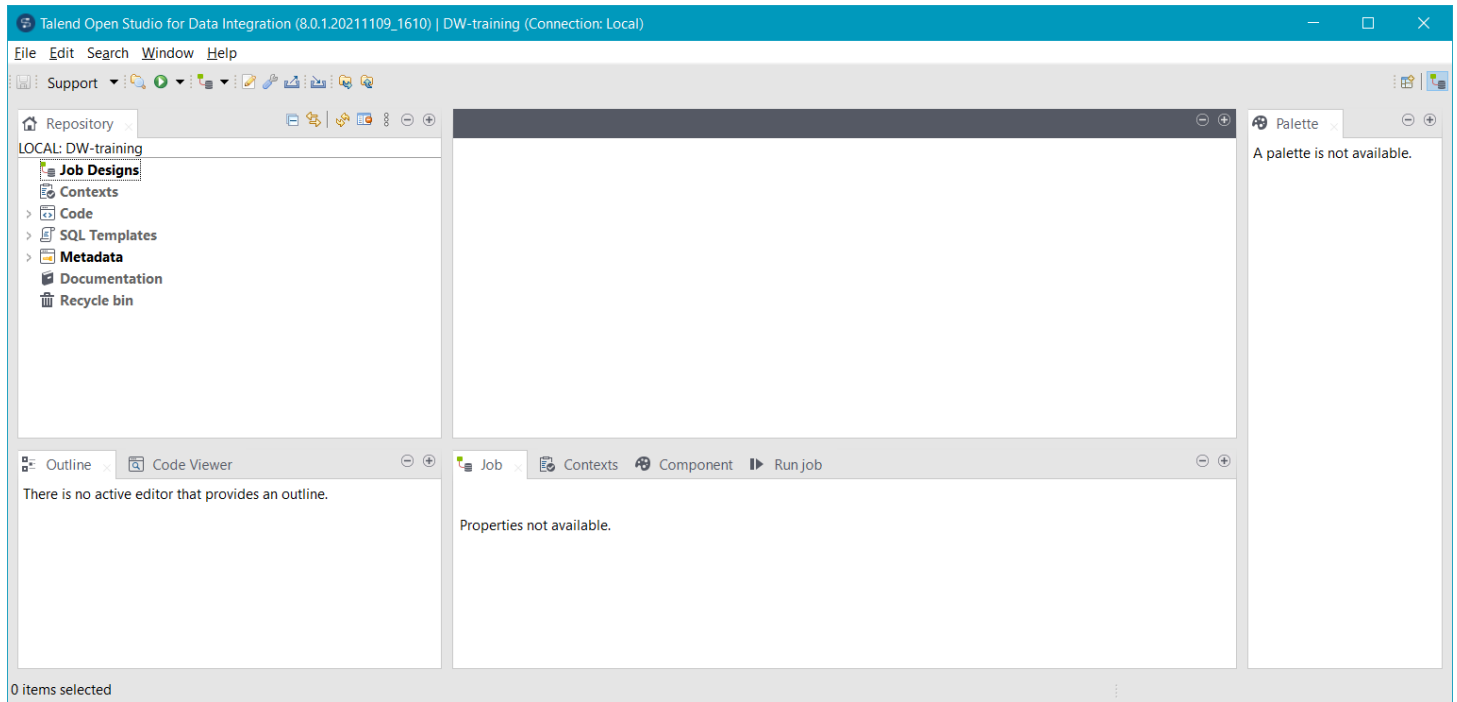
Since we are working in local (without remote ETL server) we will leave the option Local for the connection. If it's the first time we open Talend (or we want to start a new project) we should choose the option **Create a new project**, give it a name (e.g., DW-training) and press the button **Create**.

After we create a new project, it will appear in the list of existing projects (as seen in the figure below). You should select it and press **Finish** to start the tool.



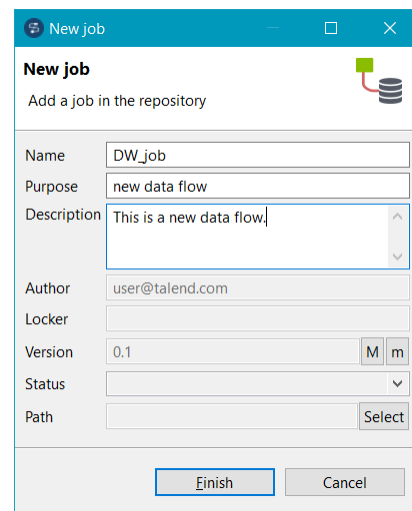
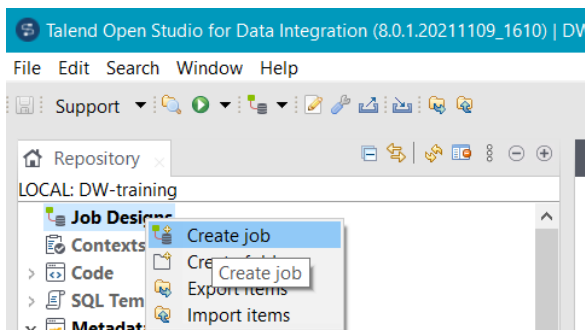
Notice that in the future, instead of creating a new project, you should choose the option **Select an existing project**, select the project you want to open and click **Finish**.

After we open a new or existing project in Talend, we are supposed to see the following starting screen (please notice that depending on the version you are running, the welcome screen can differ):

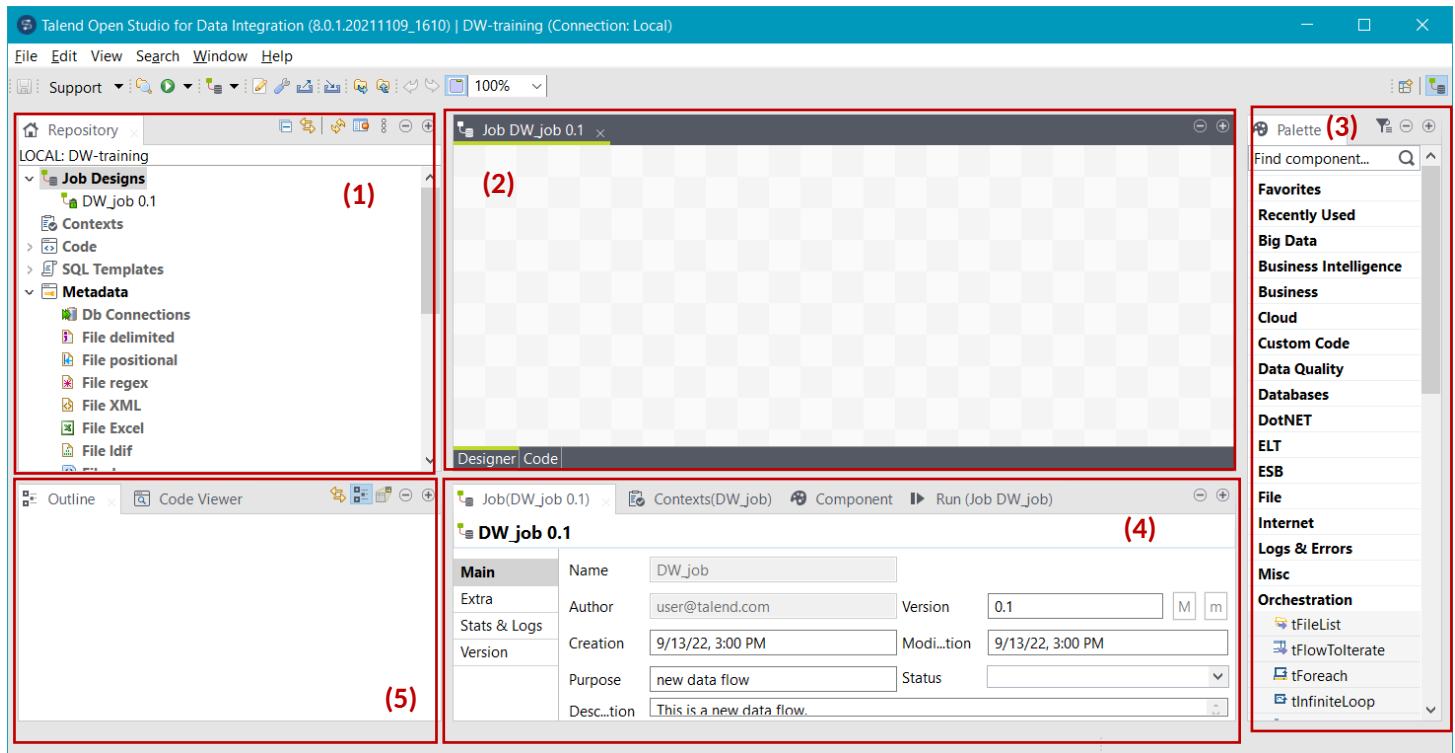


2. Creating empty data flow (job):

To create new data flow (for designing an ETL flow) we right click on **Job Designs** inside the **Repository** pane and choose the **Create job** option, as shown in the figure below-left. We should then provide at least a name for the newly created job and press **Finish**, as shown in the figure below-right.



After we create a new job, the empty design canvas should be visible in the screen as shown in the figure below.



We should distinguish 5 main panes of the Talend interface (see figure above)

- (1) **Repository** of the project where we will find all the jobs and subjobs created inside the project (under **Job Designs**), as well as the project **Metadata**, from which we will specifically use **Db Connections**.
- (2) Designer canvas of the currently selected job/subjob.
- (3) **Palette** of **components** that we can include in our data or control flows.
- (4) Control pane where we can find the **Job** tab for job and transformation configuration, as well as the **Run** tab for runtime configuration and running of the current job.
- (5) **Outline** of the ETL components included in the flow and the **Code Viewer** where the designer can see the translation of each component into Java code. For this lab, we will not be using this panel.

3. Defining data source inputs of the ETL process:

First, we need to identify the sources from where the ETL process is supposed to read the data. Considering the Flying use case, four different sources are available **1)** AMOS DB **2)** AIMS DB, **3)** maintenance-personnel-airport-lookup.csv (Personnel with the airport where they are employed), and **4)** aircraft-manufacturerinfo-lookup.csv (Aircraft with their, model and manufacturer).

Considering our example, we identify the following attributes that need to be extracted from the identified data sources:

- **Flight hours (actualArrival - actualDeparture):**
 - **actualArrival** maps to the **Flights** table of **AIMS DB** and its column **actualArrival**.
 - **actualDeparture** maps to the **Flights** table of **AIMS DB** and its column **actualDeparture**.
- **Aircraft's model** maps to the **aircraft-manufacturerinfo-lookup.csv** file and its column **aircraft_model**.
- **Year** maps to the **Flights** table of **AIMS DB** and its column **actualDeparture** with the SQL function applied to extract year from the timestamp: **extract (year from actualDeparture) as year**.
- **Number of passengers** maps to the **Flights** table of **AIMS DB** and its column **passengers**.

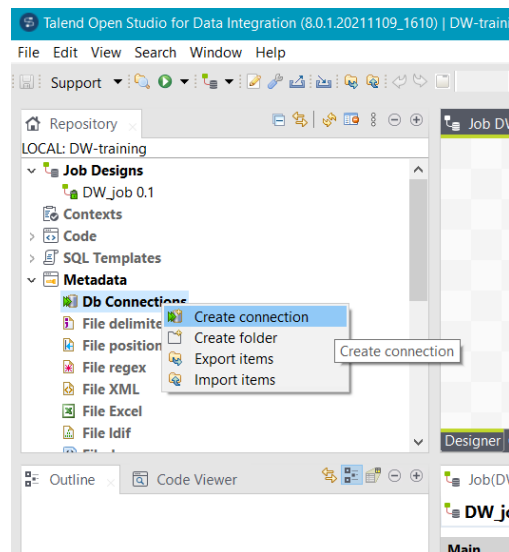
For each of these mappings (coming from a distinct data source) we need to define the data input component. Input

component templates are provided in the **Palette** pane located at the right side of the Talend window.


Defining DB connection metadata

Before we start adding data input components, we will first create a DB connection to the AIMS PostgreSQL DB, that we will use throughout this tutorial. In the similar manner, you should define other connections (e.g., with the AMOS database, or with the Oracle database for the output DW).

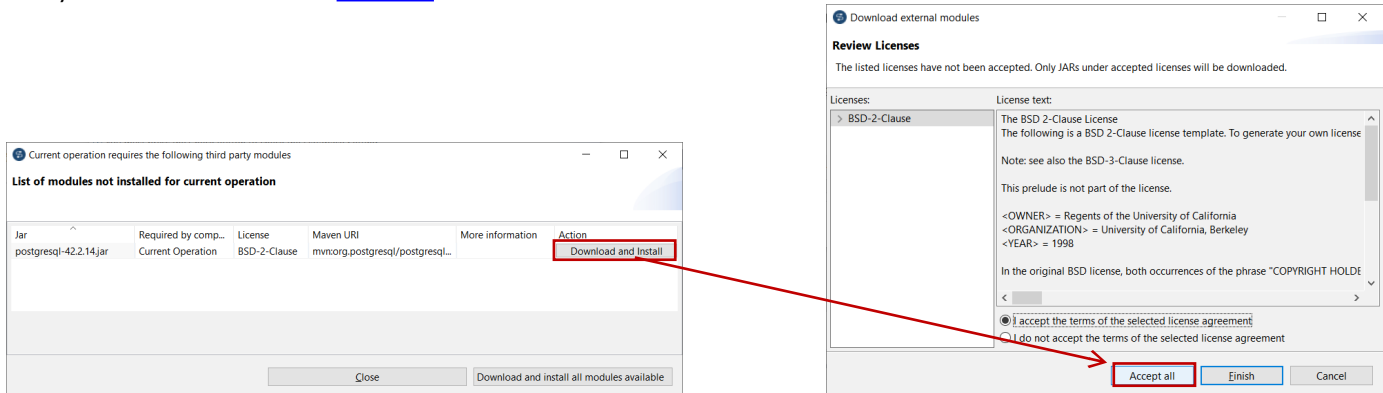
We go the **Metadata** section of the project **Repository**, right-click over the **Db Connections** option and choose **Create connection**, as shown in the figure below.



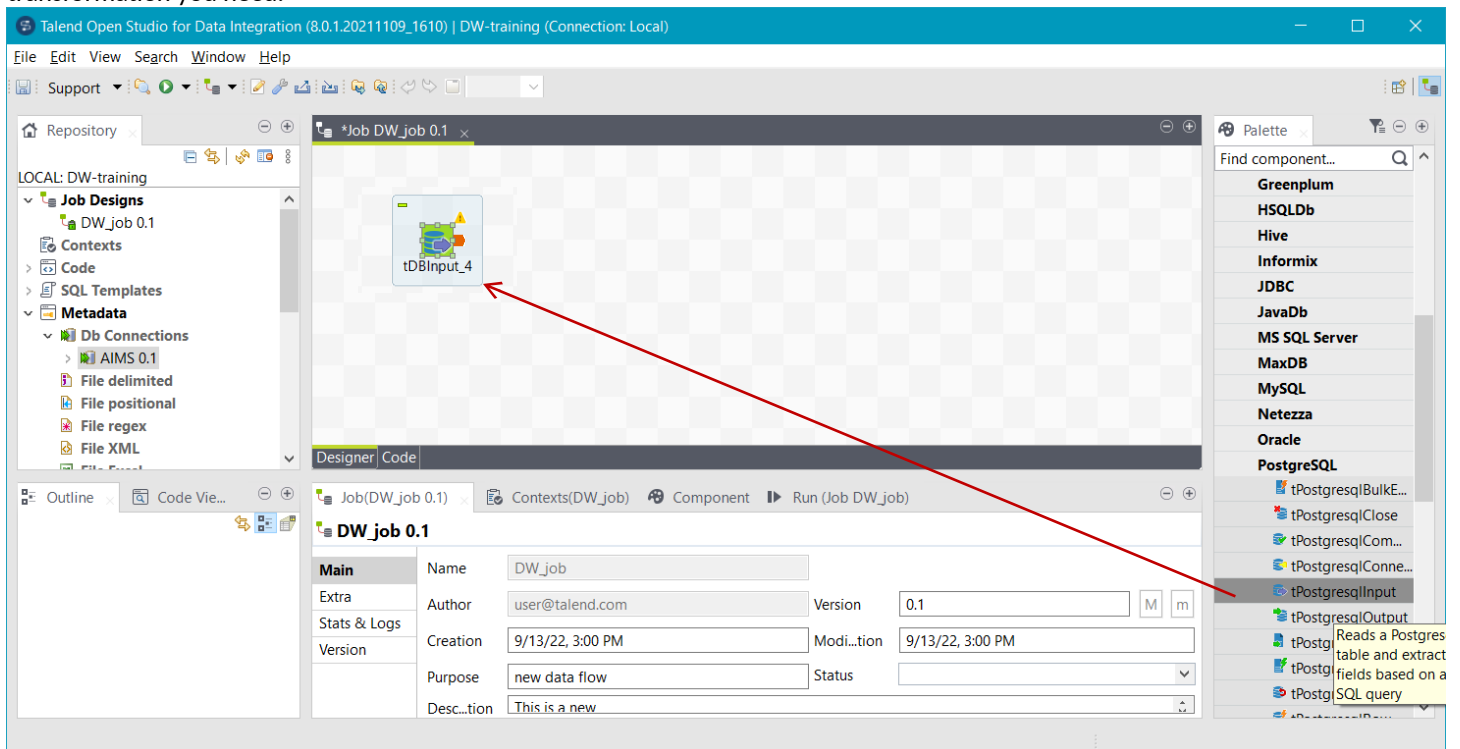
In the dialogs below, we need to give a name to the connection (e.g., AIMS), click **Next**, and then define the connection parameters as shown in the figure below (Importantly, you should use your UPC username and as a password use the provided password in the form **DBddmmyy** – where **dd** is day, **mm** is month, and **yy** is the year of your birth). Remember that you must be previously connected to the [UPC VPN](#).

A screenshot of the 'New Database Connection on repository - Step 1/2' dialog box. It contains fields for 'Name' (filled with 'AIMS'), 'Purpose', 'Description' (filled with 'DB connection to AIMS'), 'Author' (filled with 'user@talend.com'), 'Locker', 'Version' (0.1), 'Status', and 'Path'. At the bottom, there are buttons for '< Back', 'Next >', 'Finish', and 'Cancel'. A red arrow points from the 'Next >' button to the 'Test connection' button in the next dialog.A screenshot of the 'New Database Connection on repository - Step 2/2' dialog box. It shows configuration for a PostgreSQL database. Fields include 'DB Type' (PostgreSQL), 'Db Version' (v9 and later), 'String of Connection' (jdbc:postgresql://postgresfib.fib.upc.edu:6433/AIMS?), 'Login' (name.surname), 'Password' (masked with dots), 'Server' (postgresfib.fib.upc.edu), 'Port' (6433), 'DataBase' (AIMS), and 'Schema' (public). There is an 'Additional parameters' field. At the bottom right, the 'Test connection' button is highlighted with a red box. Below it are 'Export as context' and 'Revert Context' buttons. At the bottom, there are buttons for '< Back', 'Next >', 'Finish', and 'Cancel'. A red arrow points from the 'Test connection' button to the 'Check Connection' dialog.

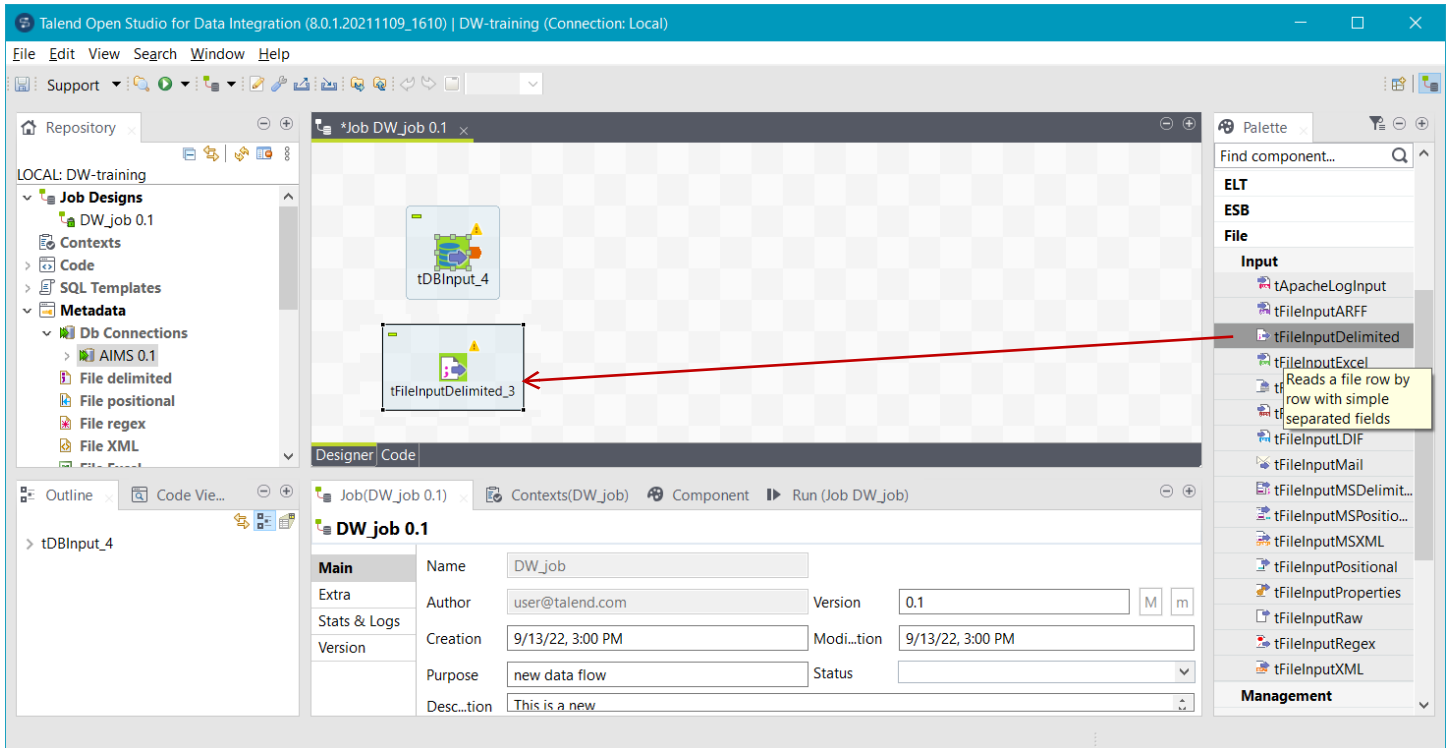
After clicking **Test connection**, only the first time, Talend will prompt a window to install a JDBC driver for connecting to a specific DBMS, in this case PostgreSQL. In the list, it will appear postgres Jar and you should click on the **Download and Install** and then **Accept all**, as it is shown in the figure below. After the installation, Talend will return the message whether the connection is successful or not, as in the above figure. If not successful, please revise the connection parameters, and ensure that you are connected to the [UPC VPN](#).



Once we have the DB connection created, we can proceed with adding a data input component to our Talend data flow. Inside the **Palette** tab, we choose the Databases->DB Specifics->PostgreSQL->**tPostgresqlInput** transformation (as shown in the figure below) and drag it to the designer canvas. Notice that you can always use the **Find component...** option to quickly locate the transformation you need.

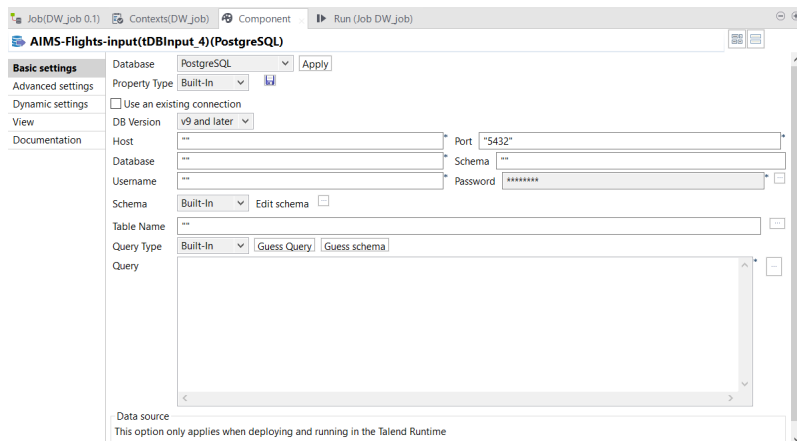


Similarly, we also add the File->Input->**tFileInputDelimited** transformation to read the data coming from the aircraft-manufacturerinfo-lookup.csv file.

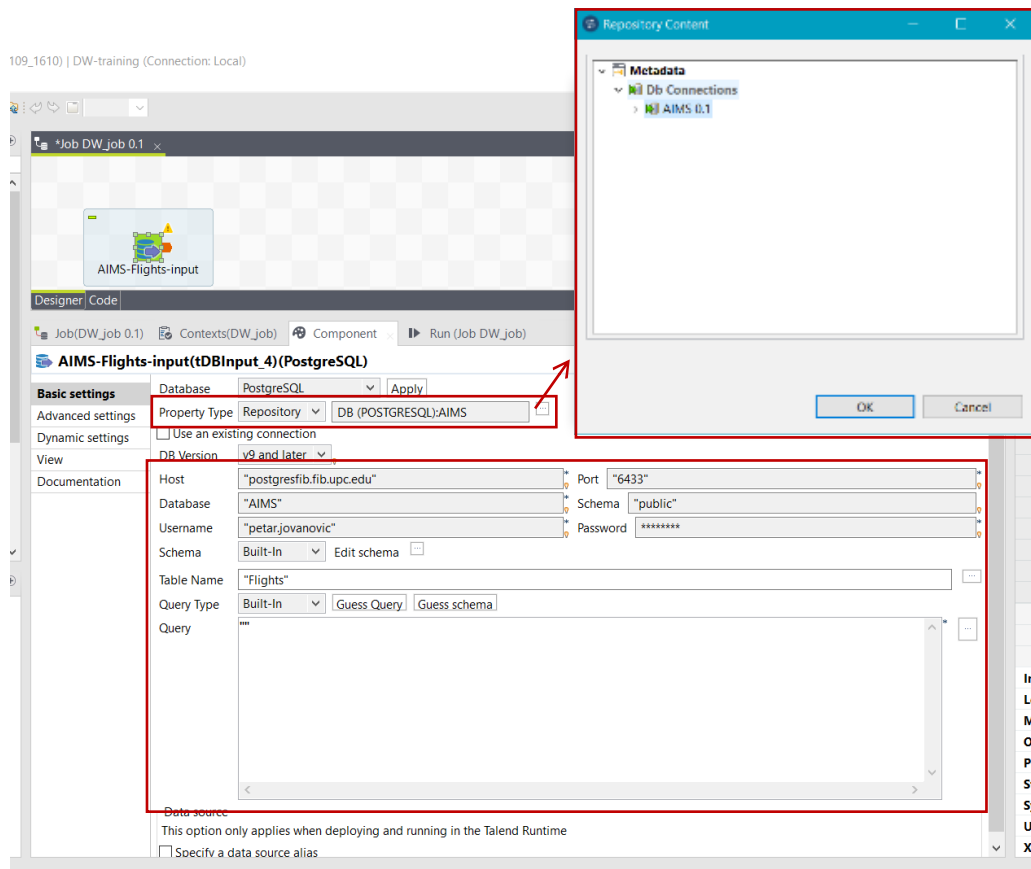


a) Parameterize table input component:

After clicking the **tPostgresqlInput** component (the inner green square, not the outer grey one), in the control pane below, under the **Component** tab, the configuration to parameterize the added component will open, as seen in the figure below.



In the **Basic settings** section we can define a new DB connection or use a previously created one. We will use the connection to the AIMS DB that we previously added to the metadata. For this, in the **Property type** section we choose the **Repository** option and then by clicking ... button open the **Repository Content** window where we can locate our Db Connection (AIMS), and click **OK**. You will see how the connection parameters below are then filled in automatically from the stored metadata.

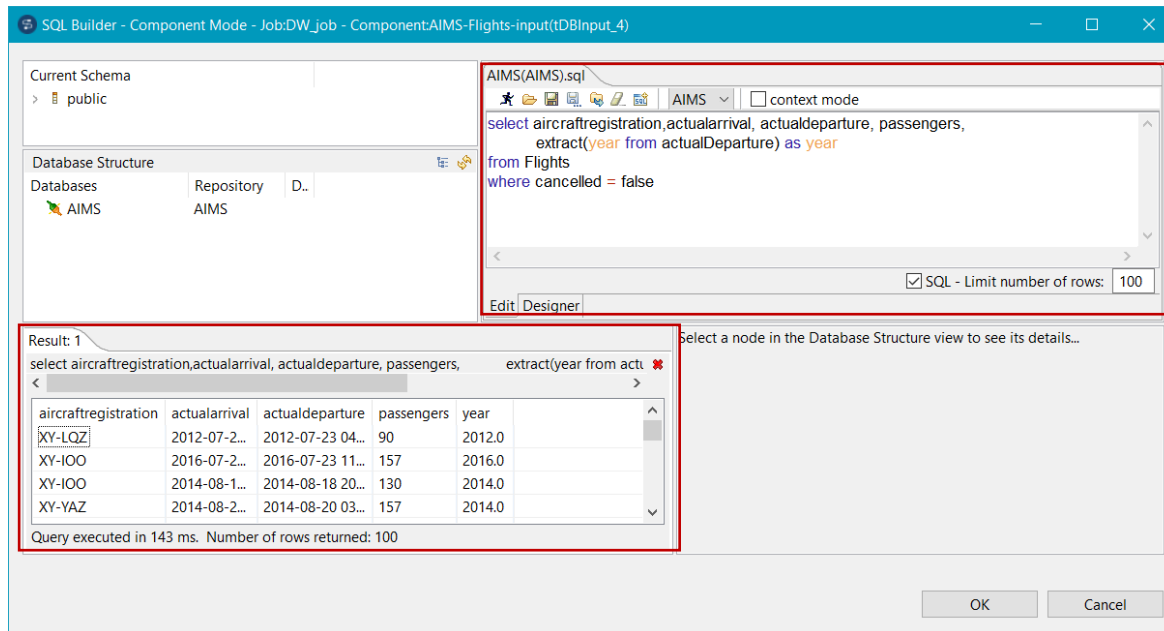


We can as well define the table name from the given DB (in this case we write “Flights”).

The next step is to obtain the required columns from the given table. For that in the **Query** section we click on ... and open the SQL Builder window.

There we can write the SQL query³ to obtain only the needed data from the table Flights. Notice that since we are calculating the flight hours, we are interested only in the flights that were not cancelled, thus we also add a WHERE clause with the condition “cancelled = false” to the SQL query. We can execute the given query and visualize the sample rows from the table (100 by default), by clicking the running human button. We can use SQL Builder to test several queries and see what kind of data it is giving us. Once we are satisfied with the query to extract the data, we can click OK (See figure below).

³ As we do not need the data of all the columns to be read we can specify only those needed (avoid using * to get the whole table) and/or apply operations to extract parts of other attributes (e.g., extract (year from actualDeparture)).

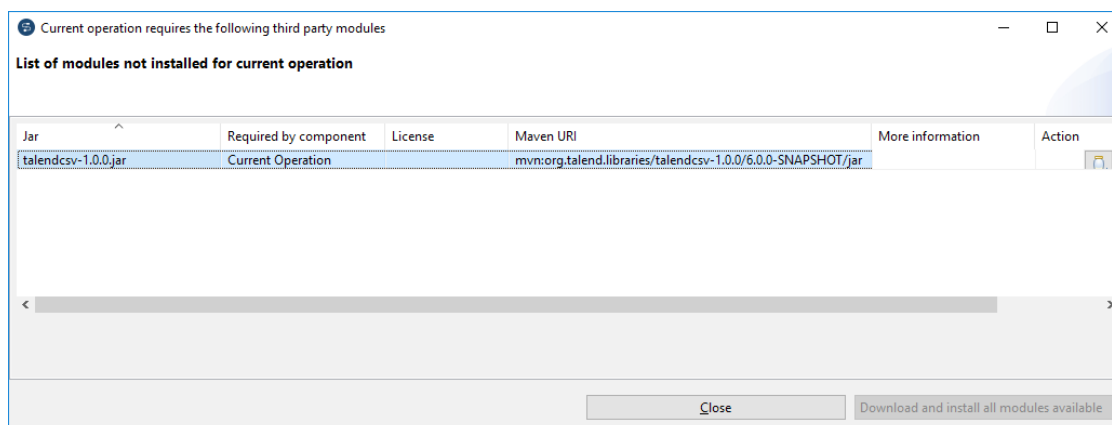


We can see then that the provided query is placed inside the **Query** section of the component configuration. After this, in order to define the output schema, automatically from the given query, you should click on the **Guess schema** option that will automatically list the output columns from the query with data types *guessed* from executing a sample of the input data.

Importantly, notice that this step would require an external 3rd party jar library (**talendcsv**) to be installed.

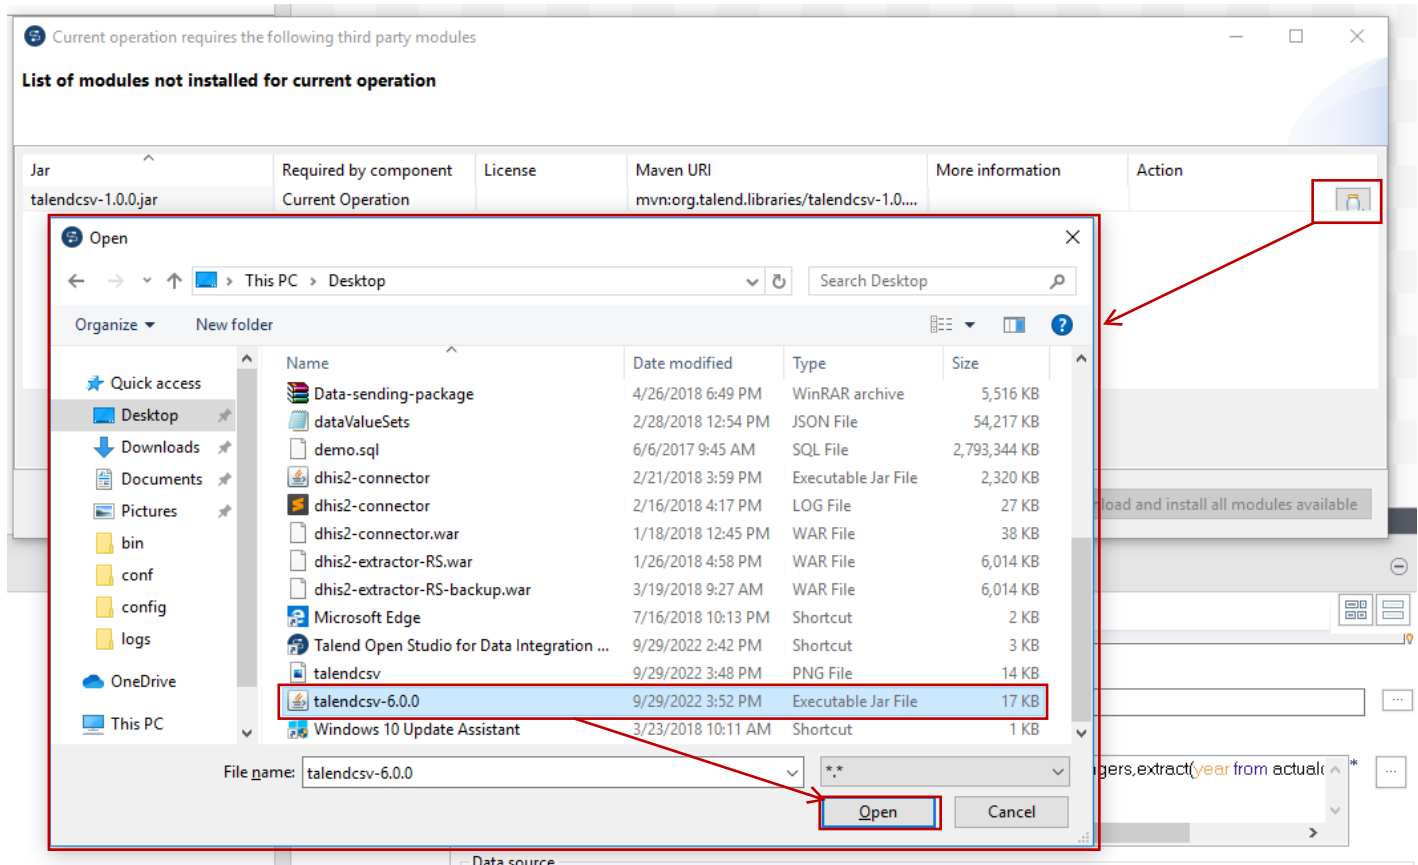
Installing 3rd party libraries in Talend

In some cases, if an external library is missing in Talend, and it is not included in the installation package, it must be downloaded and installed manually. Here we will show this process for the **talendcsv** library, used for handling csv files as well as obtaining schema information automatically. If you click on the **Guess schema** option like in the previous step, you will initially get the following window :



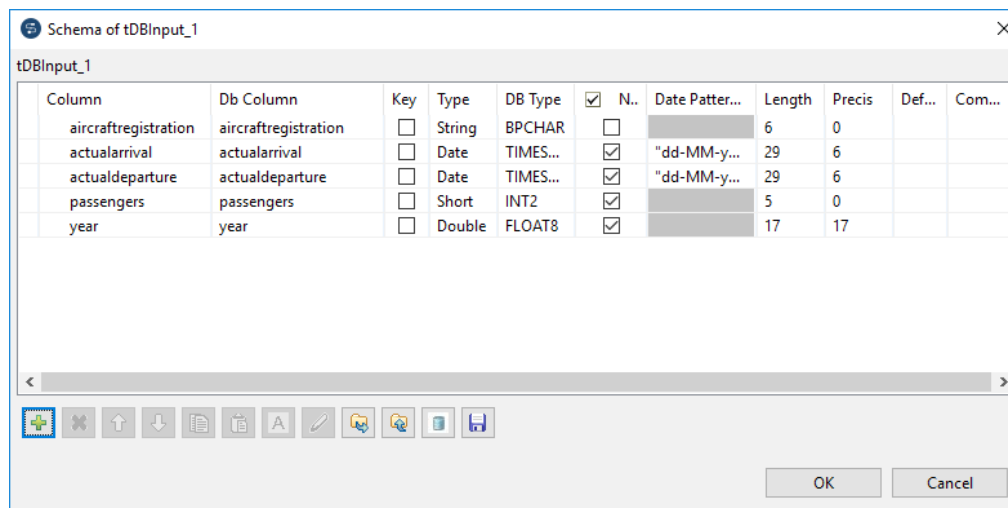
Notice that in this case, there is no option to download and install this module automatically, like in the case of PostgreSQL JDBC library. In such case, we must search for the external library in the online [maven repository](https://artifacts-oss.talend.com/nexus/content/repositories/TalendOpenSourceRelease/org/talend/libraries/talendcsv/6.0.0/talendcsv-6.0.0.jar)⁴, download it and then point to it from Talend, as shown in the figure below.

⁴ <https://artifacts-oss.talend.com/nexus/content/repositories/TalendOpenSourceRelease/org/talend/libraries/talendcsv/6.0.0/talendcsv-6.0.0.jar>



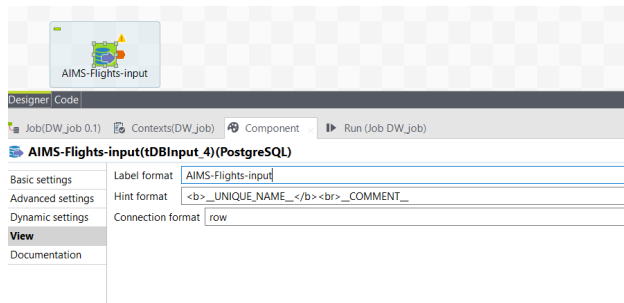
Once the 3rd party library is successfully installed in Talend, you should be able to use the given functionality.

In this case, if we click again the **Guess Schema** option, you should be able to see the automatically *guessed* schema with columns and their datatypes, as seen in the figure below.



You can edit it, if necessary. Click OK and it will store the schema metadata for this component.

In addition, as part of the **Component** configuration tab, under the **View** section (**Label format**) we can also change the name of the component and give it a more project-specific name. See the figure below.



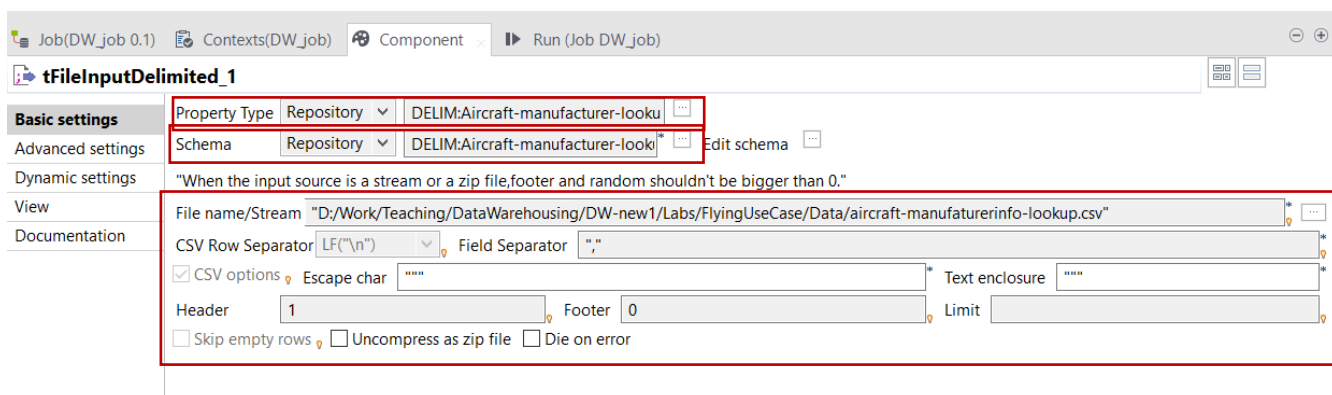
Important note: Notice that various operations that are to be performed inside the ETL flow can be as well supported by the source engine (in this case PostgreSQL DBMS). For example, filtering, various format conversions, operations over temporal attributes (date/time difference, interval overlapping, etc.). Therefore, for the sake of optimizing data processing, and to avoid loading more data than needed into an ETL engine, we can push those operations to the source engine. In that case, they should be implemented within the SQL query as shown above.

b) Parameterize **tFileInputDelimited**⁵ component:

Defining CSV file/schema metadata

Just like with the DB connection, Talend allows us to centralize the definition of the metadata needed to connect and read CSV files (i.e., delimited file). Before you start parametrizing the **tFileInputDelimited** component, please follow carefully the steps explained in the Talend web page⁶ to define the connection to the Aircraft manufacturer info lookup file.

- (1) As with any other component, we can define the name of the step by entering the desired name in the **Label format** text box inside the **View** section of the **Component** tab.
- (2) In the **Basic settings** section of the **Component** tab we can then define the configuration to read the data from the CSV file. Since we previously created a centralized metadata to treat this CSV file, we should simply select this metadata under **Property Type** and **Schema** sections from the **Repository**. Once we do this, the configuration and schema will be filled in automatically. Alternatively, we can use **Built-in Property type** and **Schema**, and create a one-time configuration for this component. The advantage of using the centralized metadata is that we can easily reuse them various times inside the same project. See the figure below.



4. Adding the data transformation components of the ETL process:

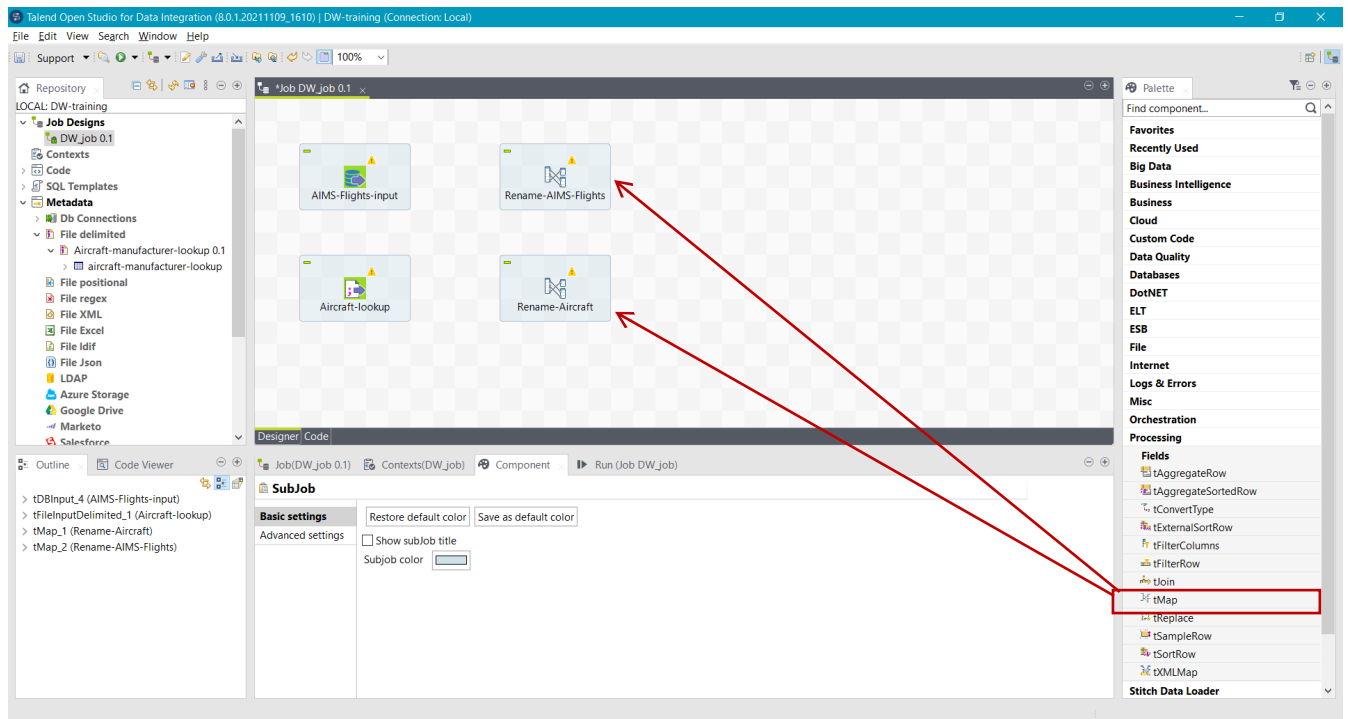
After we defined the input data components to extract data we need to define the transformation steps of the ETL process.

- (1) First, for each input data step we can add a *rename* step that will assure that all the field names in

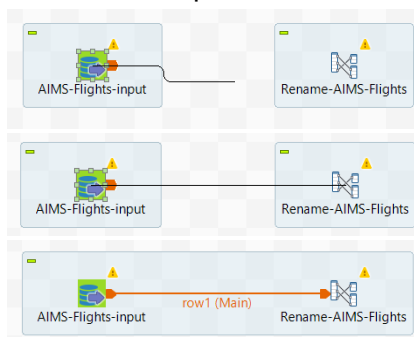
⁵ <https://help.talend.com/r/en-US/8.0/delimited/tfileinputdelimited>

⁶ <https://help.talend.com/r/en-US/8.0/studio-user-guide-esb/centralizing-file-delimited-metadata>

the process are unique. This step can be important as various independent sources may have the same names for the fields with different semantics. From the **Palette** we choose the Processing->**tMap**⁷ component and drag it to the designer canvas. See figure below.

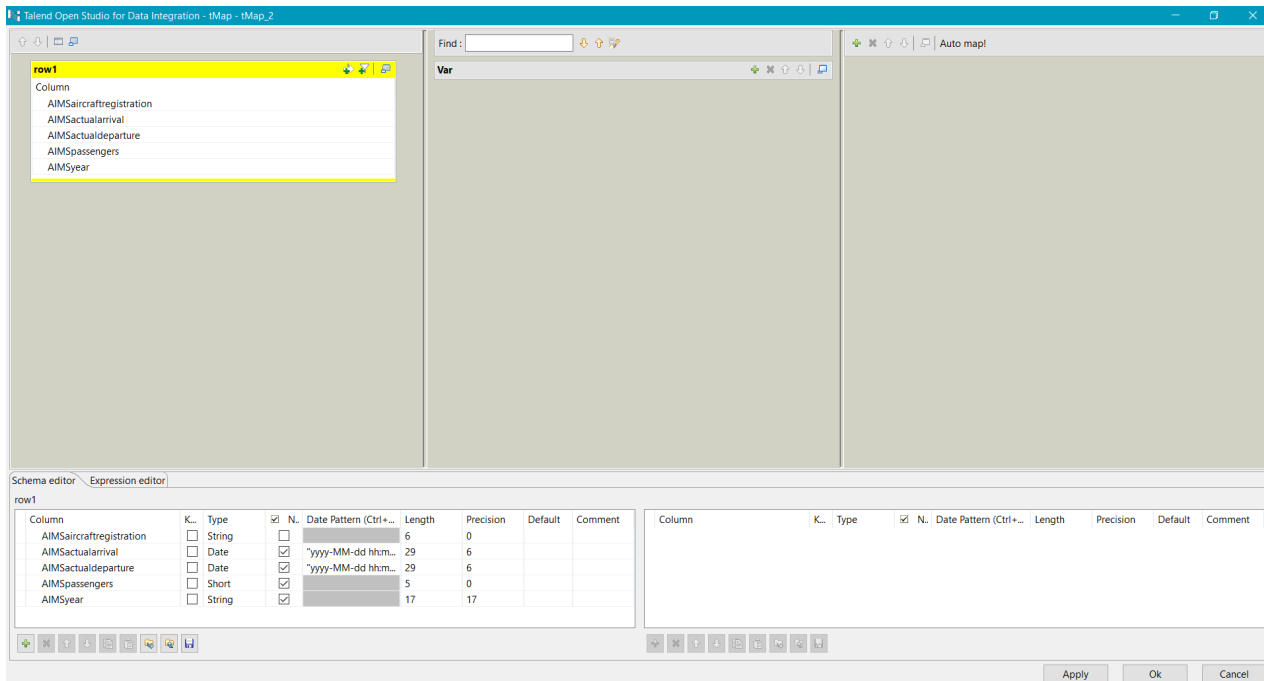


- (2) To insert the edge from the table input or csv file input components to the rename component (or between any other two components in the canvas), we click on the first component and then a small red arrow will appear on its right (). When we click and hold the red arrow it will activate edge creation and we just need to drag it towards the component to which we want to connect. Once we release the mouse button it will establish the data flow edge between the two components. See the figures below. Another interesting consequence of this action is that these two components will become part of the same *subjob* (*sub data flow*).

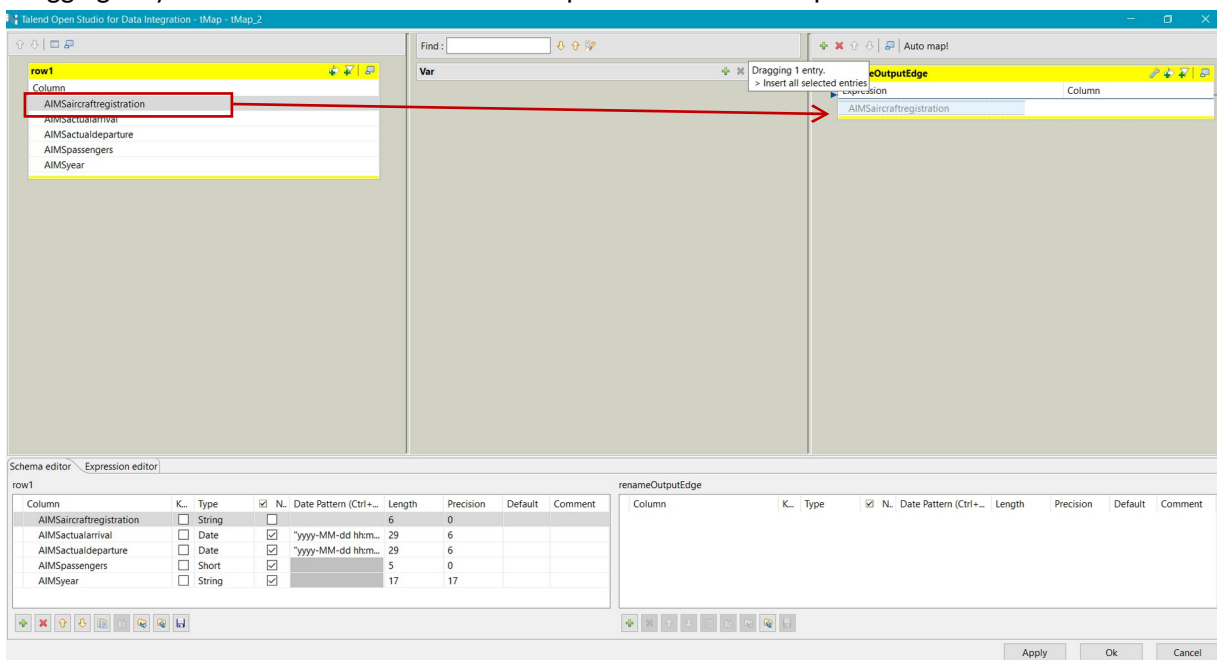


- (3) To parameterize the **tMap** component for renaming the columns, we need to double click the **tMap** component which opens the **tMap** configuration window as seen in the figure below. Then, inside the **Schema editor** tab below we simply need to click on the column we want to rename and give it another name. For example, to distinguish the columns coming from different sources, we can add as a prefix the name of the source (e.g., AIMS). Besides renaming inside the same Schema editor we can also alter the metadata of input fields (i.e., data type, length, precision, or format of the date). After we finish, we click **Apply** and **OK** to save the configuration.


⁷ tMap component has very rich semantics: For more details refer to: <https://help.talend.com/r/en-US/8.0/tmap/tmap-component>



Later on, after we connect the tMap component to the next component in the data flow (i.e., after completing step 4) below), we need to come back and map the input schema to the output schema, by dragging only those attributes that we want to pass to the next component.



In this case we map all 5 columns to the output



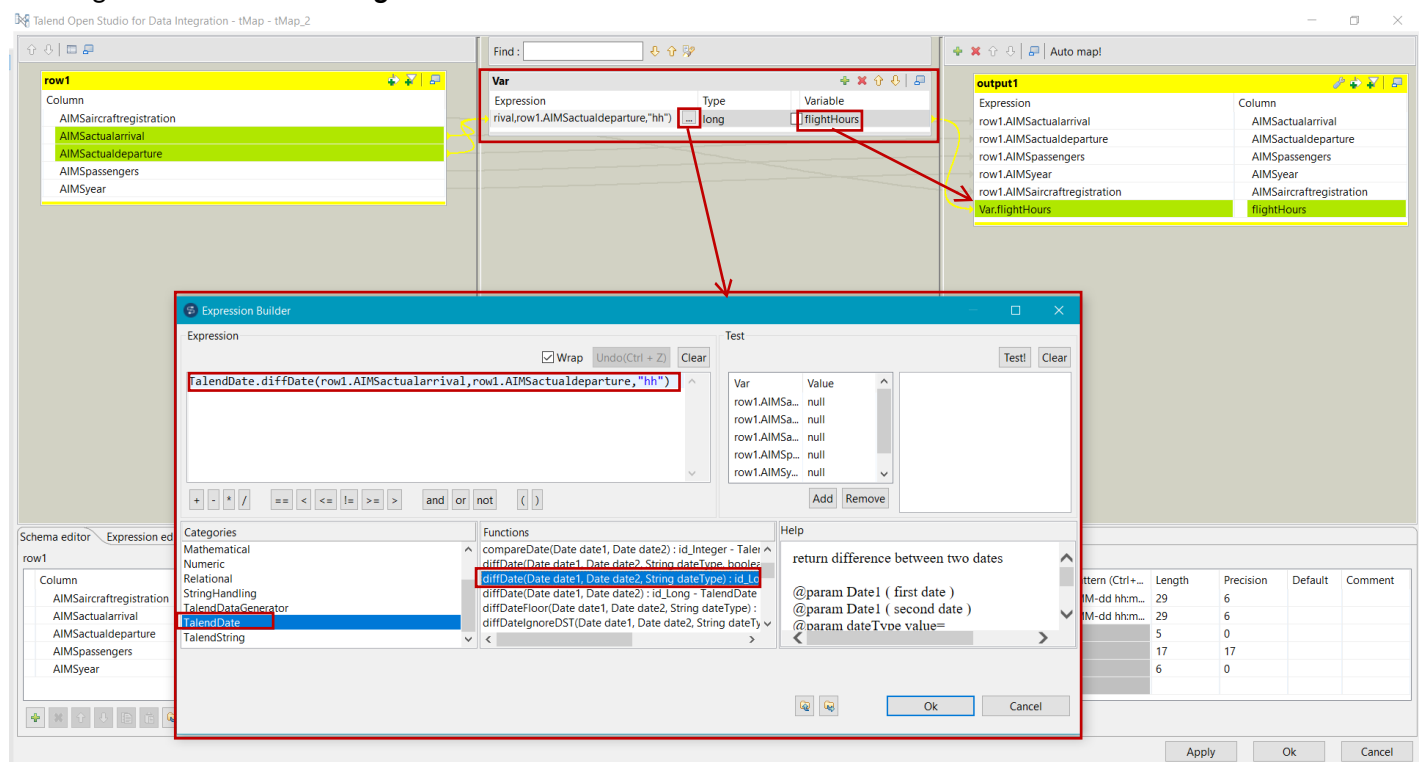
Notice that with this functionality we can do projections of the columns, omitting the columns that are no longer needed in the flow.

Notice that this way we performed renaming at the input of the tMap component (i.e., in the left hand side

Schema editor), which means that the renamed columns will be propagated to any output in the rest of the flow. Similarly, instead of doing it at the input, we could as well do the renaming at the output of tMap (i.e., in the right hand side **Schema editor**), but only after we connect the tMap component to another component(s) in the ETL flow and have the columns properly mapped to that output. This way, the renaming should be done for each output of the tMap (in case there is more than one), and the columns can get different names different outputs.

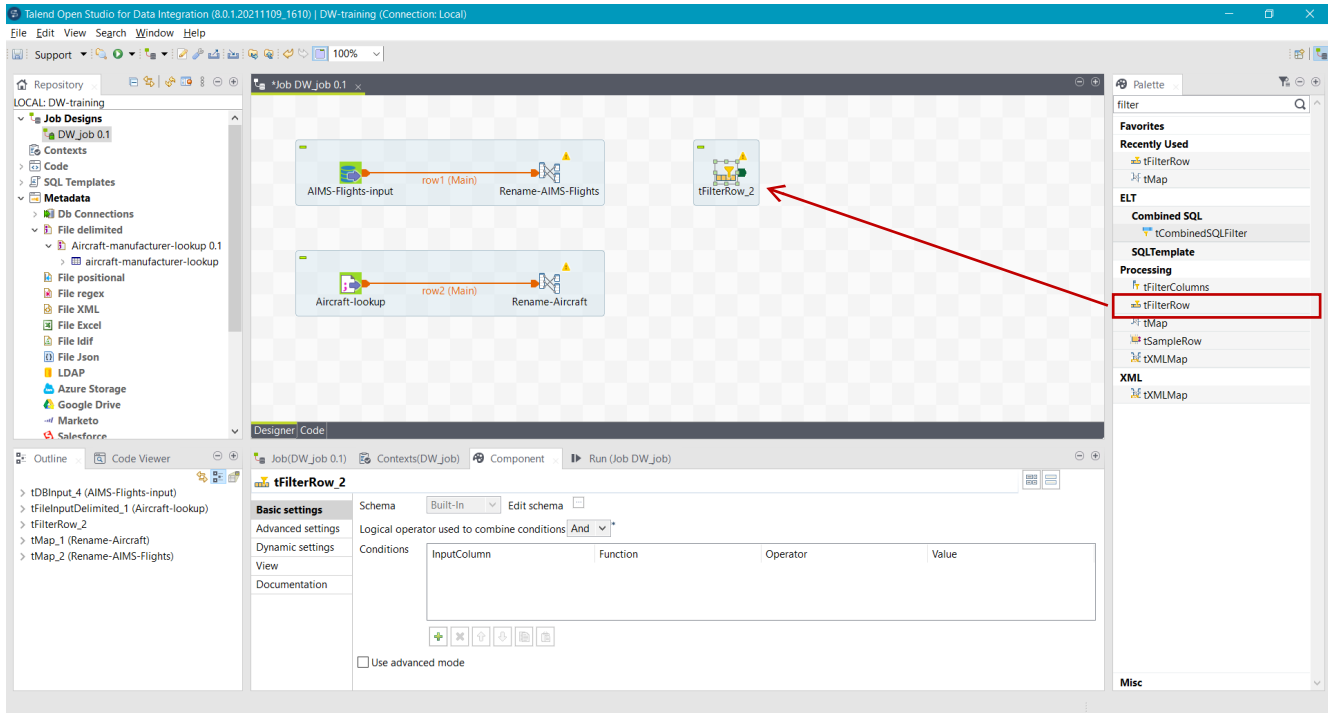
The tMap component has very rich semantics and we will use it more throughout this tutorial. For more details, refer to: <https://help.talend.com/r/en-US/8.0/tmap/tmap-component>

Indeed, we will take advantage and derive a new variable with an expression to calculate the flightHours attribute as required in KPI 1. Inside the same configuration we will add the expression (made out of input attributes and define the new variable, which we will then pass to the output. See the figure below. Inside the Expression Builder window, in this case, we will use a special Talend operator for dealing with the dates (**TalendDate**). More precisely to calculate the difference (**diffDate**) between the timestamp of the arrival and the timestamp of the departure in terms of hours ("hh"), in order to calculate the flight hours of each flight. Notice that this way, we use the same **tMap** component to do the renaming of the columns and deriving the new attribute - **flightHours**.

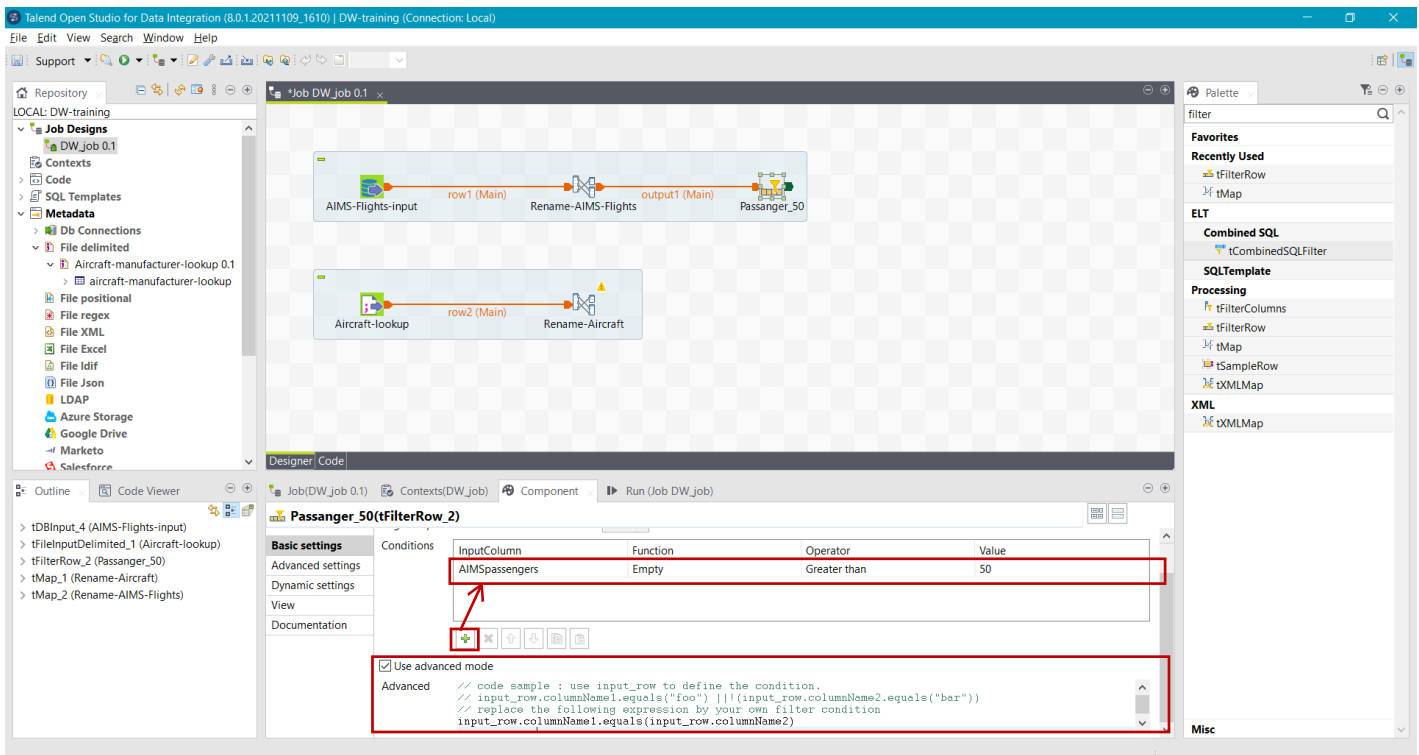


(4) Then, we need to add a *filtering* step for the restriction defined in the KPI 1(i.e., *passengers > 50*).

From the **Processing** category of the **Palette** choose **Processing->tFilterRow** component and drag and drop it to the designer canvas as shown in the figure. Before parameterizing it, we need to connect it to the previous **Rename-AIMS-Flights** step and configure the **tMap** as explained previously.

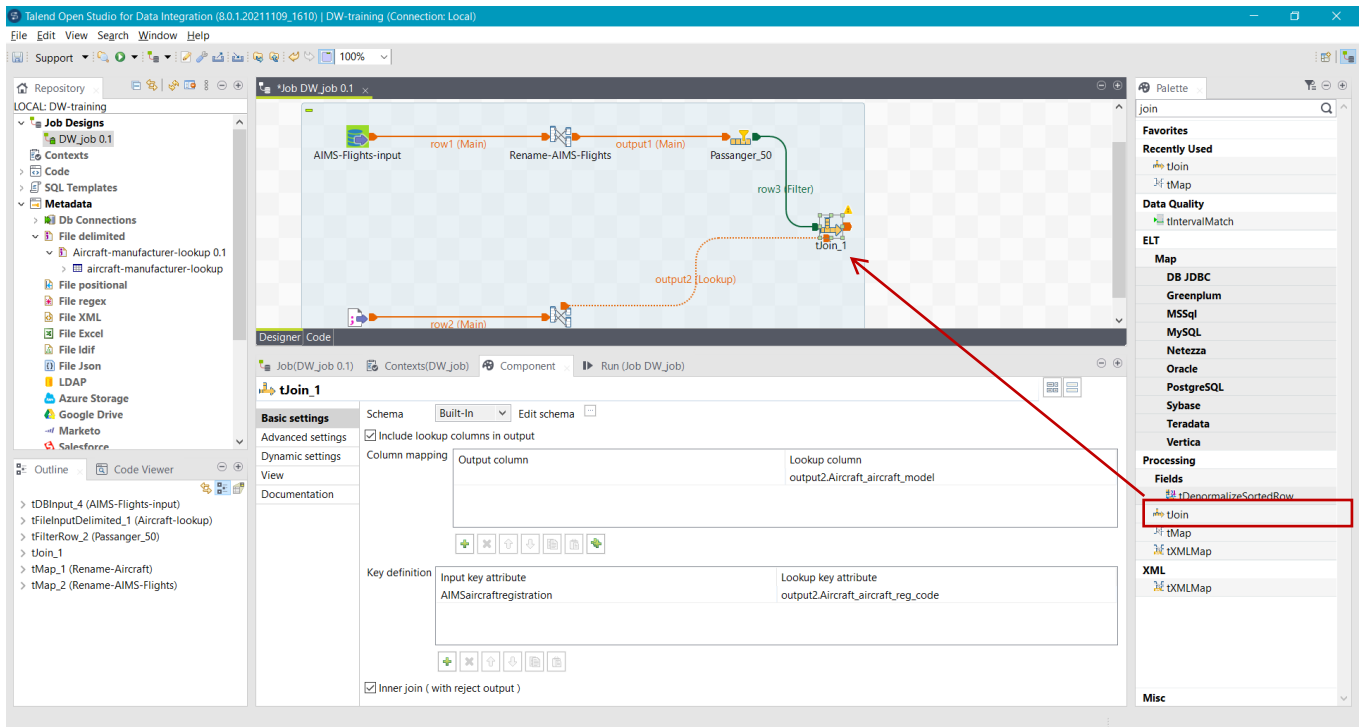


(5) We then need to parameterize the **tFilterRow** component considering the restriction provided in the input KPI 1. See figure below.

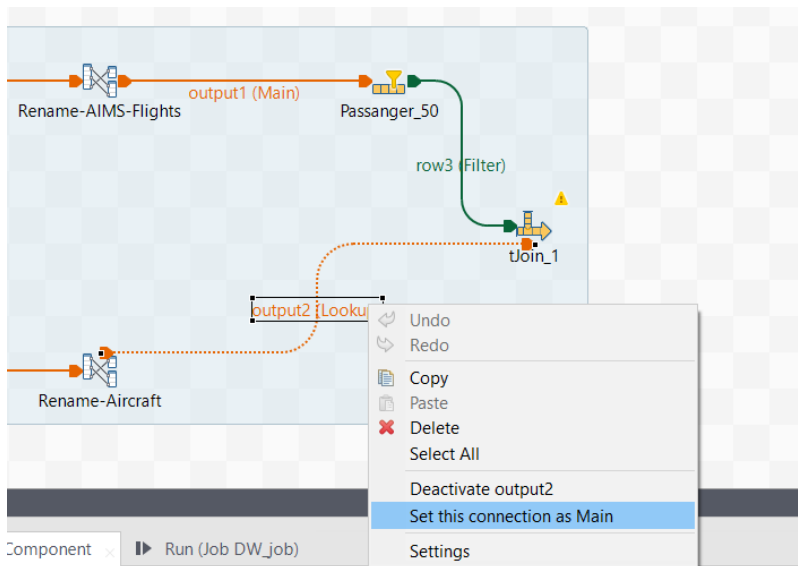


Important note: while **tFilterRow** provides you with typical predefined predicates to filter the values of input fields, to implement more complex Filtering operations, you can use the advanced mode and write your filters in Java. In this case, we will leave it unselected.

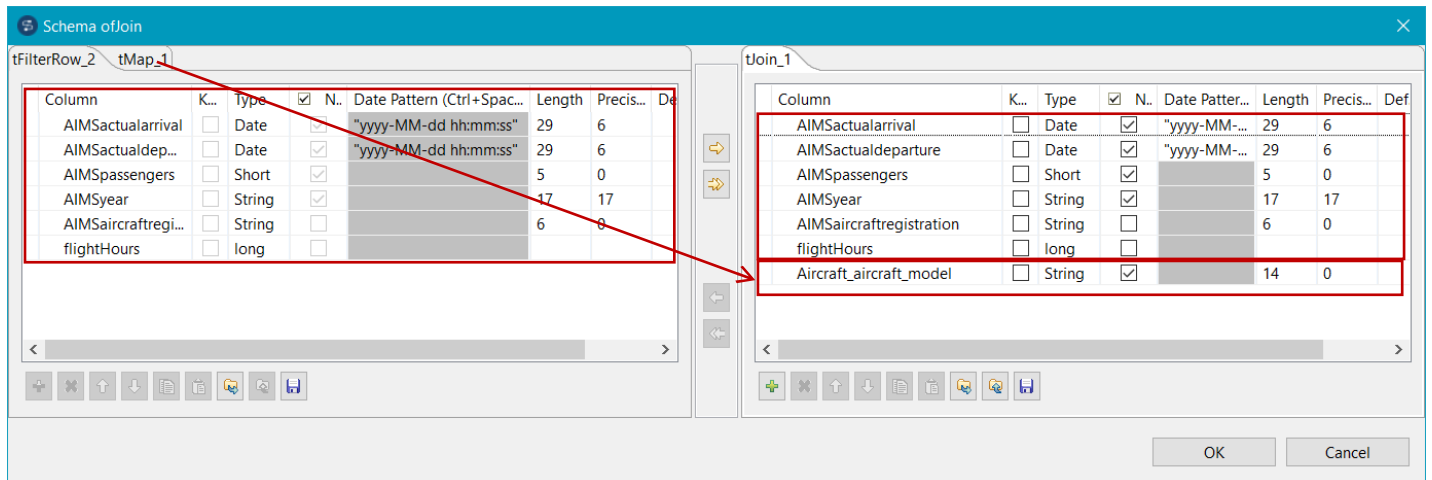
- (6) Then we add the *Join operations* to relate data coming from the two defined data sources (i.e., AIMS Flights table and Manufacturer lookup table from the CSV file). From the **Processing** category of the **Palette** we choose the **tJoin** component and drag it to the designer canvas as shown in the figure below. As before, we need to connect it to the previous steps in the flow and in addition configure **tMap** of the **Rename-Aircraft** step such that the correct columns are mapped to the **output2**.



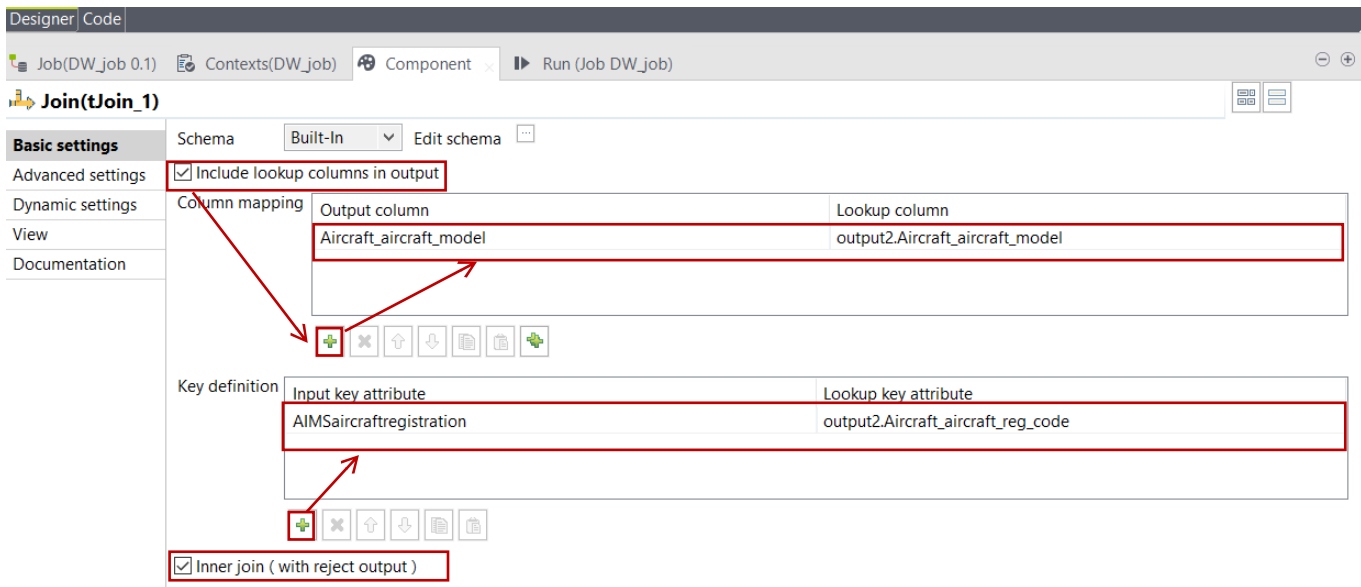
- (7) Then we parameterize the tJoin step to properly relate in this case rows coming AIMS Flight table (**Main** connection) and rows coming from Aircraft Manufacturer Lookup CSV file (**Lookup** connection). By default, the first connected step is **Main** and the second **Lookup**, but we can change this a posteriori by right-clicking on the connection and choosing the other option (see the figure below). Notice that this serves to optimize the execution of the join, and that the larger data input should be **Main** which is read as a stream, while the Lookup is a smaller and cached in memory.



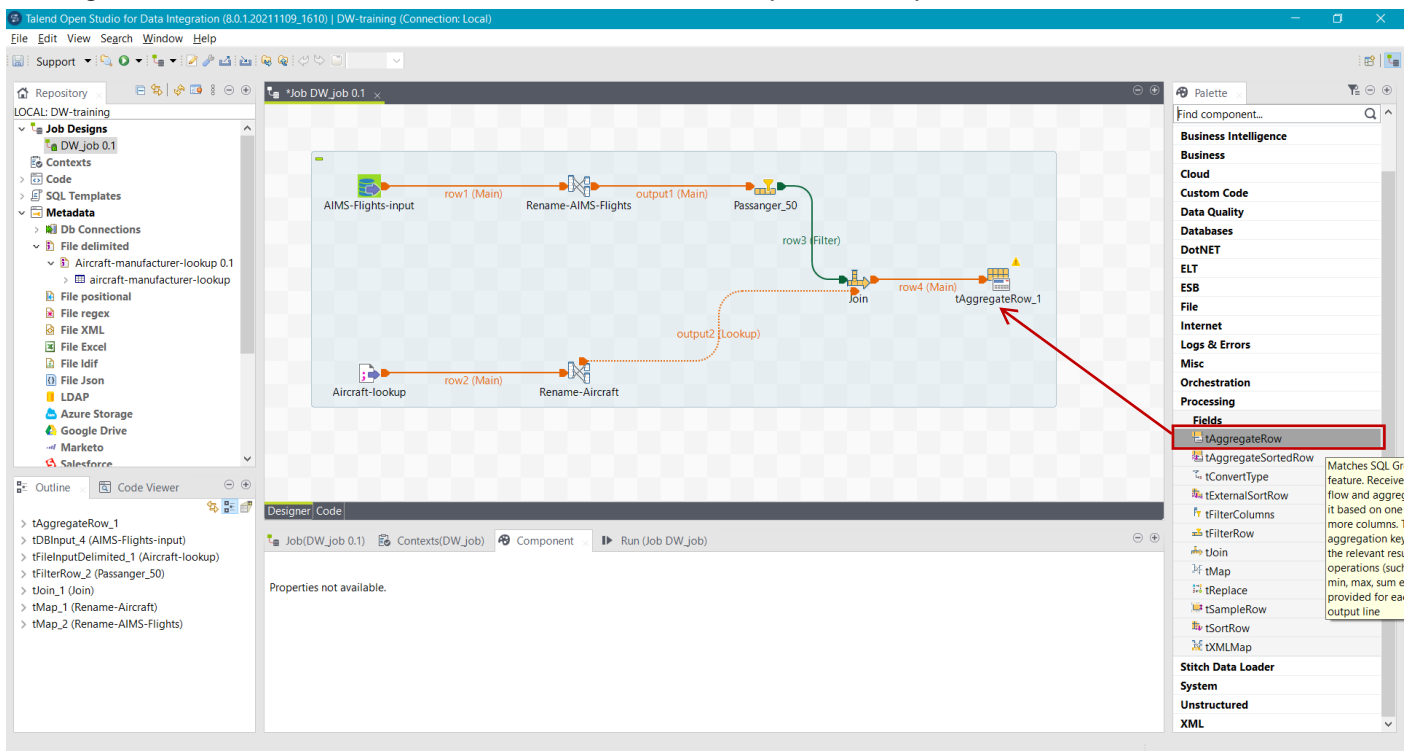
We then need to configure what will be the output schema from this join step. Clicking in the Component tab on **Edit schema ...**, opens a window where we can define which attributes from the Main and from the Lookup input step we want to pass on. On the left we see two tabs, one for each input step schema, from which we can select the attributes and pass them to the right under the join output schema. See figure below. We pass all the attributes from the Main input step and only the **Aircraft_aircraft_model** from the Lookup input step.



Lastly, define the join attributes (**Key definition** section), as well as include the lookup attributes that we want to pass to the output (**Include lookup columns in the output** checkbox, and **Column mapping** section). See the figure below. In this case, as defined in the schema, we only need aircraft model in the output, but we need to additionally add the mapping to the Lookup column. In addition, we also choose to use the Inner join (**Inner join (with reject output)** checkbox), such that the Flights for which we do not have an entry in the lookup file are not passed to the main output. If this option is not selected, the join will act as an *outer join* and by setting Main and Lookup connection we can make it LEFT or RIGHT outer join.



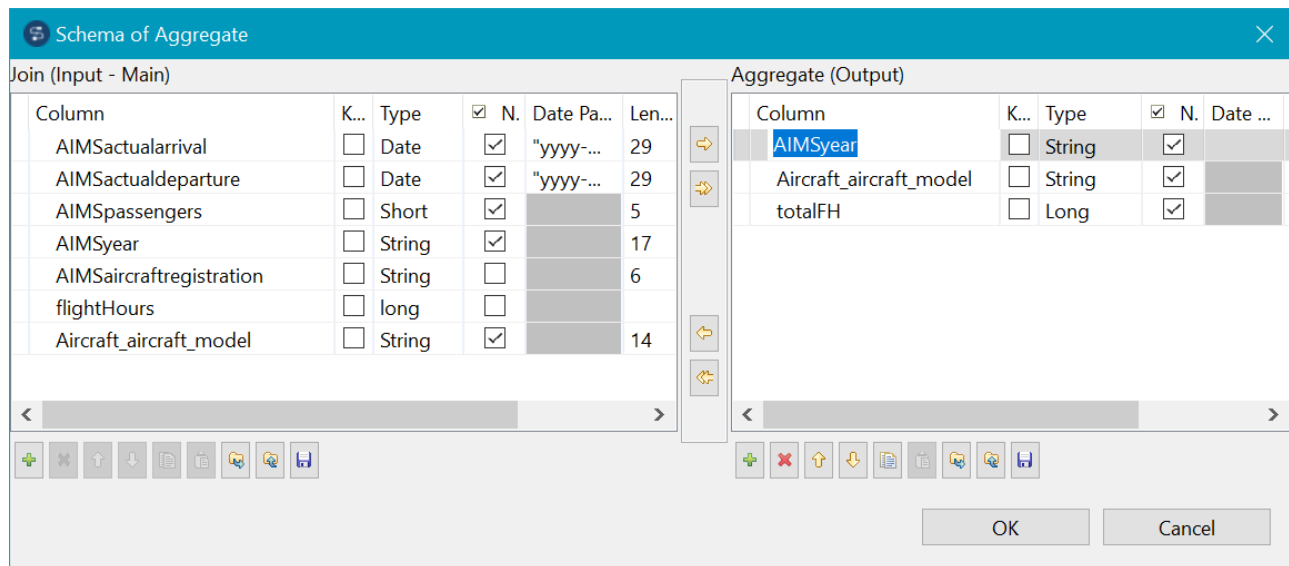
- (8) Afterwards, we need to aggregate data considering the given KPI 1. From the **Processing** category of the **Palette** we choose the **tAggregateRow**⁸ component and drag it to the designer canvas as shown in the figure below. Afterwards we also need to connect it to the previous step in the data flow - **Join**.



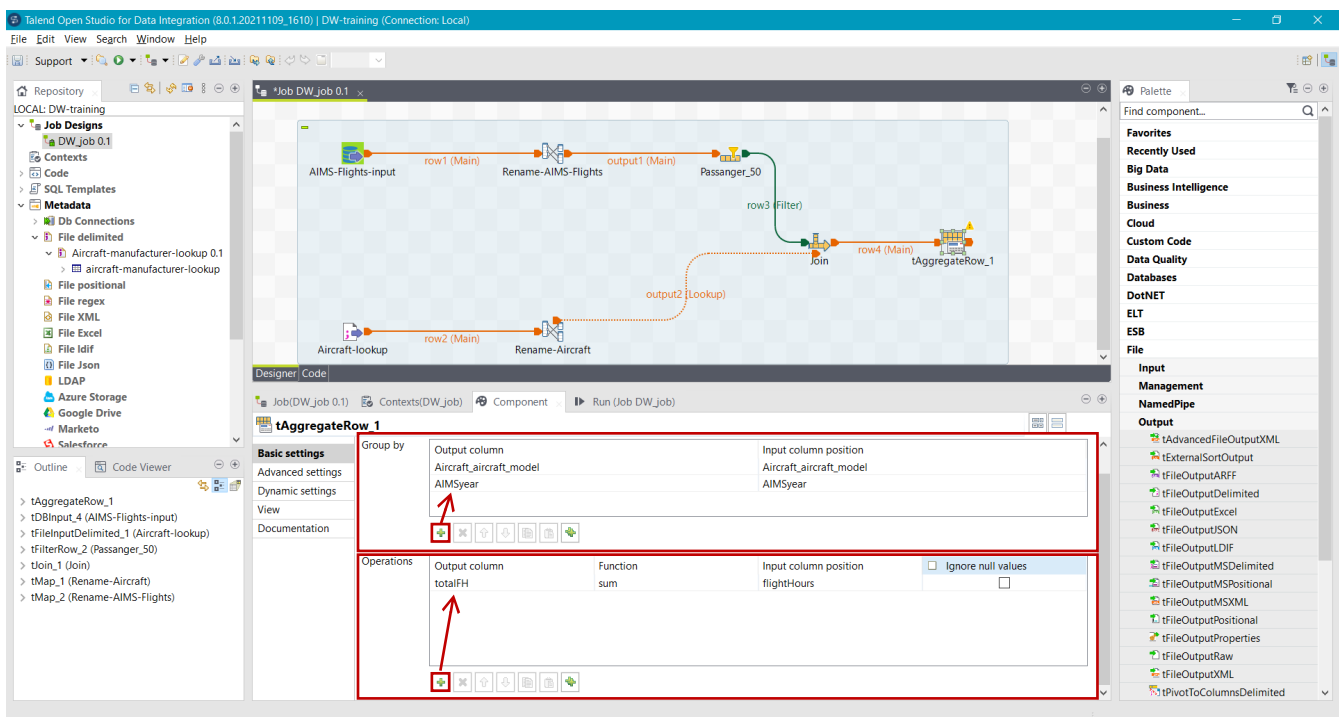
- (9) We then parameterize the **tAggregateRow** component considering KPI 1. First of all, like in the case of **Join**, we need to define the output schema of the operations, by clicking **Edit schema ...** In the schema definition window, we need to select the attributes/columns from the input that we want to send to the output. These would later serve as group by attributes (**Aircraft_aircraft_model** and **AIMSyear**). We then need also to add a new output

⁸ Notice that there is also **tAggregateSortedRow**, which requires that the input rows are sorted on the Group By keys, and can optimize the execution if the input rows already follow such order.

column that will serve to receive the value resulting from the aggregation operation (**totalFH** and choose the **long** as a data type). See figure below.



(10) We then finally configure the aggregation component by specifying the **Group by** columns and the aggregation **Operation** (on which input column to be applied – **flightHours**, function – **sum**, and which output column to fill in with the result – **totalFH**). See figure below.



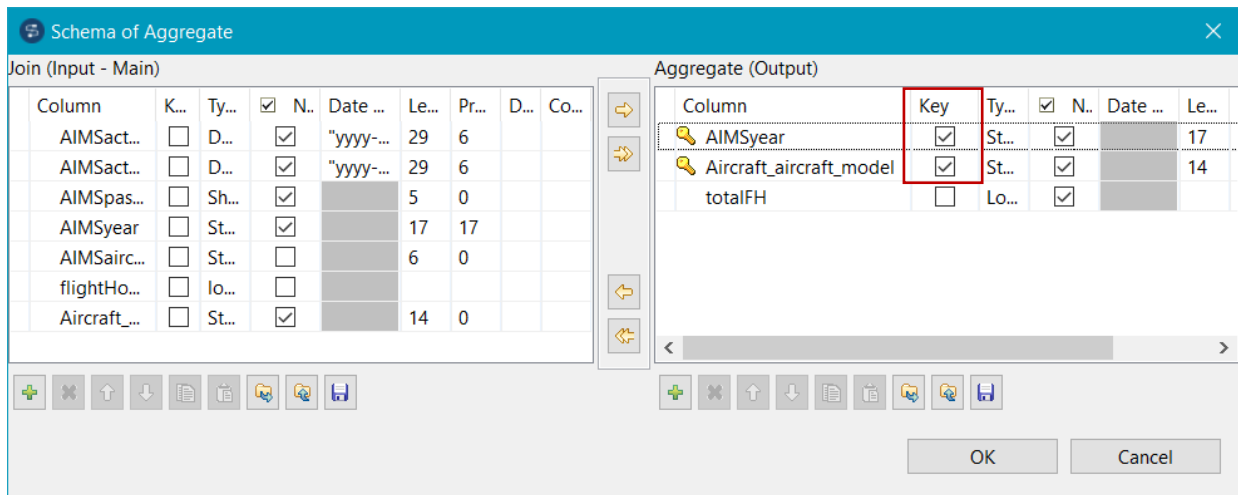
(11) Finally, after we finish the data flow, we complete it by loading the data resulting from the data flow to the output DW table. In this case, our DW is in Oracle database, so we need to add **tOracleOutput** component from the **Palette (Databases -> DB Specifics -> Oracle -> tOracleOutput)**. To do this, we need to previously add the connection to our Oracle database to the Metadata repository (See the steps in

Defining DB connection metadata section and parameters in the figure below). Like in the previous case, first time you add an oracle connection, you may be asked to install the corresponding JDBC driver.

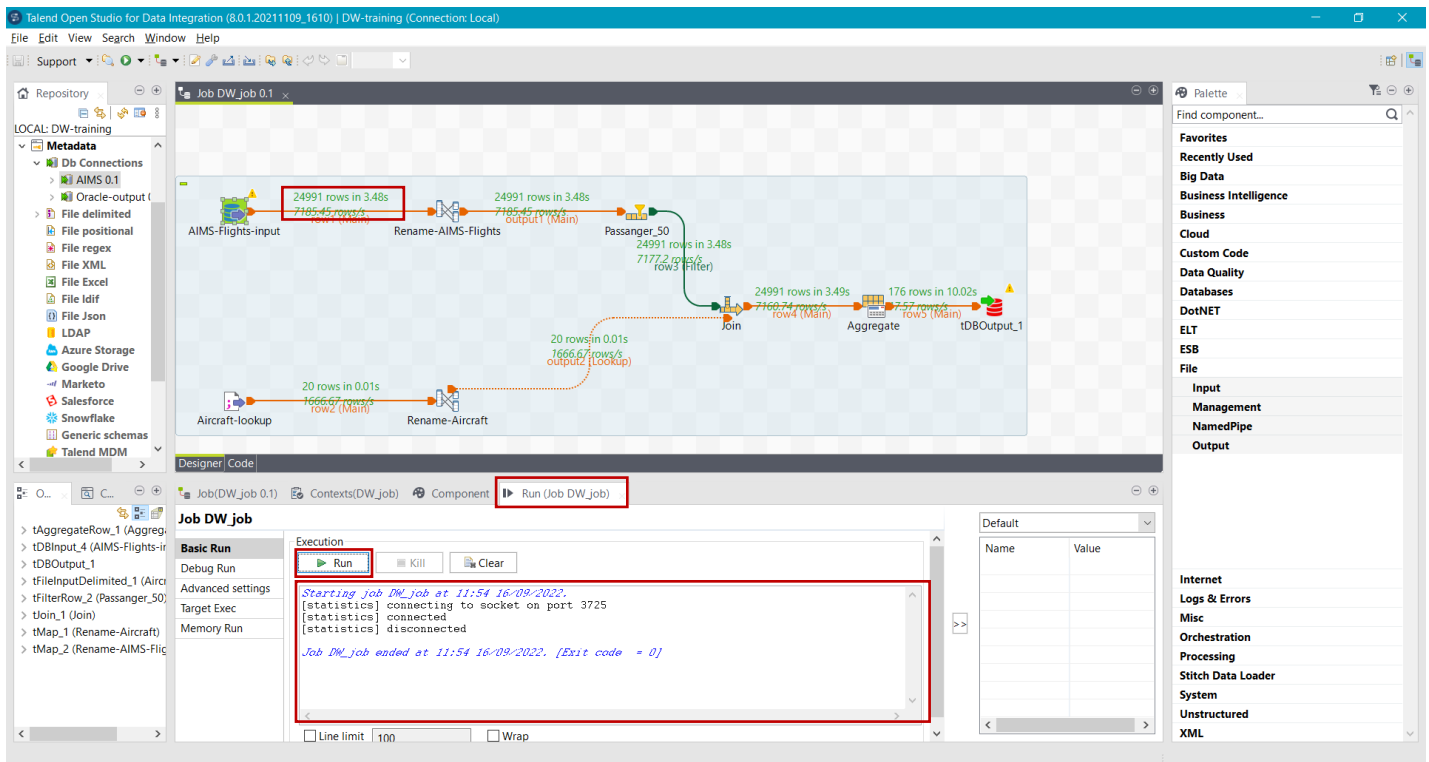
We then configure the **tOracleOutput** component as in the following figure. Notice that we choose the previously created connection from the **Repository** as the **Property Type**.

We add the name of the output table and we select the actions **Create table if does not exists** and **Insert or update**, to (re)create a KPI1 table where data will be stored and to insert new data or update existing data in the KPI1 table. Notice that for the option **Insert or update**, we need to have clearly defined keys of the table we are inserting. This can be done either directly in the database by setting a PRIMARY KEY, or in the same data flow by editing the output schema of the previous aggregate step (**Edit schema ...**) and

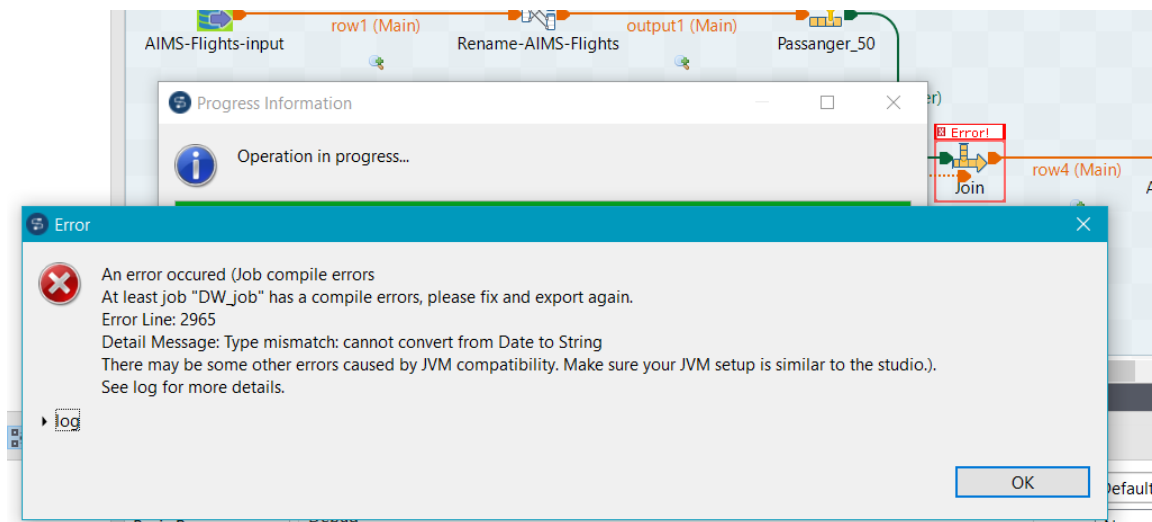
setting the **Key** property for the group by columns. See figure below.



- (12) After we create the ETL process that answers the given KPI 1 we can run the flow by opening the **Run** tab of the control panel. Under the **Basic Run** section, we can just click **Run** and wait for the job to finish. During the execution we will see some quick statistics about the number of rows and latency of the flow in each of the connection. In addition, below we will also find a log that will inform about the flow execution and possible error that may occur. In this case, there are no errors and the job finishes with success (i.e., Exit code = 0)



If there are any **compile time errors** in the created ETL flow, upon running, you should get an error message indicating what could be the issue, while the component where the error occurred will be circled in red with **Error!** message above the icon, like in the figure below (e.g., because of the type mismatch).



However, an error can also occur at **runtime**, for instance, due to faulty database connection or because of bad data at input.

Debugging runtime errors in Talend

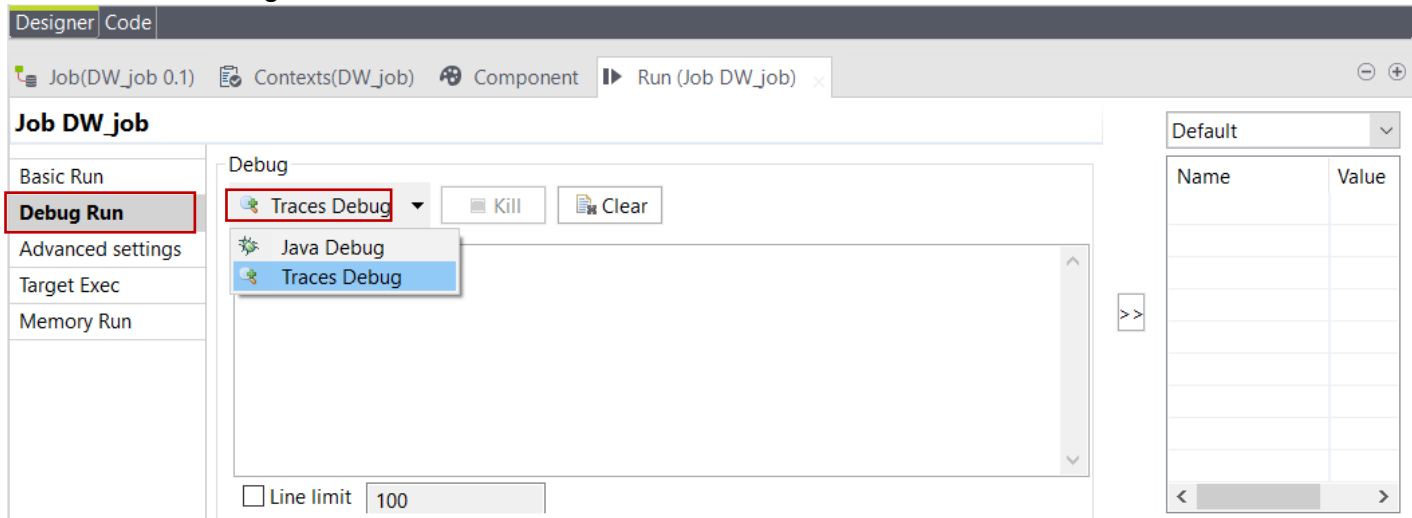
In the case, some error occur in the execution (e.g., due to the faulty data like NULL values), Talend provides tools to conveniently debug your ETL flow – The **Traces Debug** mode.

In the Traces Debug mode, Talend Studio provides the capability to monitor the data row-by-row as it flows between different components in an ETL flow. You can choose to **activate** or **deactivate** Traces or decide what processed columns to display in the traces table that displays on the design workspace when launching the current Job. You can either choose to monitor the whole data processing or monitor the data processing row-by-row or at a certain

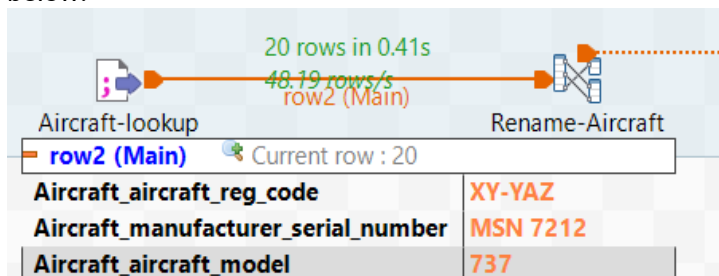
breakpoint.

To debug the Job in the Traces Debug mode:

1. With your Job open in Talend Studio, open the **Run** tab in the control panel and then select **Debug Run**.
2. Click the **Traces Debug** button⁹.

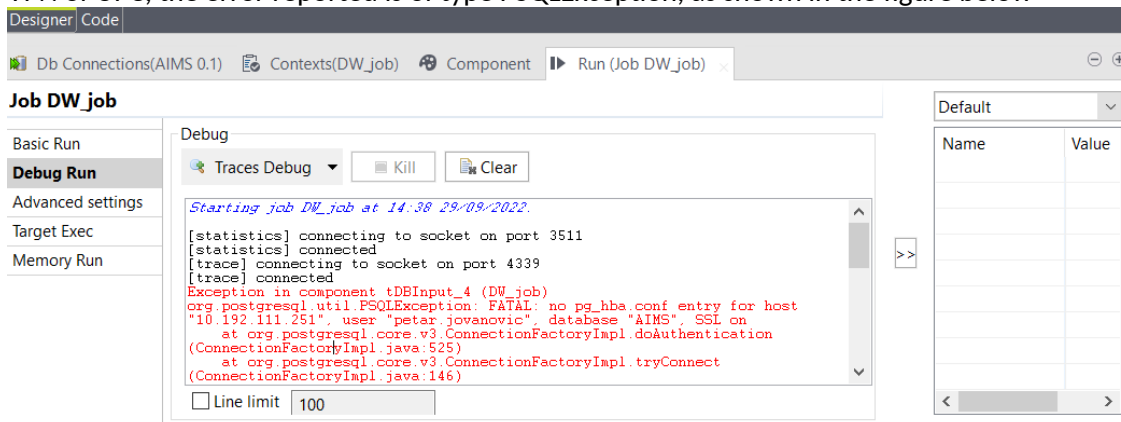


The Job runs, with row content displayed under each main output of a component, like it is shown in the figure below.



You can watch the content of the rows as they flow through the ETL flow until it encounters an error.

3. If an error occurs because of a not properly configured database connection, or if you are not connected to the VPN of UPC, the error reported is of type `SQLException`, as shown in the figure below



4. However, if an error occurs because of bad data (which is more often), you can examine the values of each field in the row that caused the error. In this case, the error reported is a null pointer exception or error when inserting a NULL value to an output database table, as shown below.

⁹ Note: Click the triangle in the right part of the button and select Traces Debug if Traces Debug does not appear on the button.

The screenshot displays the Talend Studio interface for a job named 'Job DW_job 0.1'. The top section shows a data flow diagram with components like 'AIMS-Flights-input', 'Rename-AIMS-Flights', 'Passanger_50', 'AIMSaircraftregistration', 'AIMSactualarrival', 'AIMSactualdeparture', 'AIMSpasengers', 'Aircraft-lookup', 'Rename-Aircraft', 'Join', 'Aggregate', and 'tDBOutput_1'. Data flow statistics are shown above the components, such as '1000 rows in 49.3s' and '20.28 rows/s'. The bottom section shows the 'Job DW_job' configuration with 'Debug Run' selected. The 'Debug' tab is active, showing a console with the following output:

```
Starting job DW_job at 15:15
29/09/2022.
[statistics] connecting to socket on port 3458
[statistics] connected
[trace] connecting to socket on port 4613
[trace] connected
ORA-01400: cannot insert NULL into
("ABEURP19"."KPI1"."AIRCRAFT_AIRCRAFT_MODEL")
ORA-01400: cannot insert NULL into
```

The 'row4 (Main)' data preview is also visible, showing the following data:

| row4 (Main) | Current row : 13 |
|--------------------------|------------------------------|
| AIMSactualarrival | Mon Jan 28 16:10:50 CET 2013 |
| AIMSactualdeparture | Mon Jan 28 13:44:26 CET 2013 |
| AIMSpasengers | 149 |
| AIMSyear | 2013 |
| AIMSaircraftregistration | XY-YAZ |
| flightHours | 2 |
| Aircraft_aircraft_model | null |

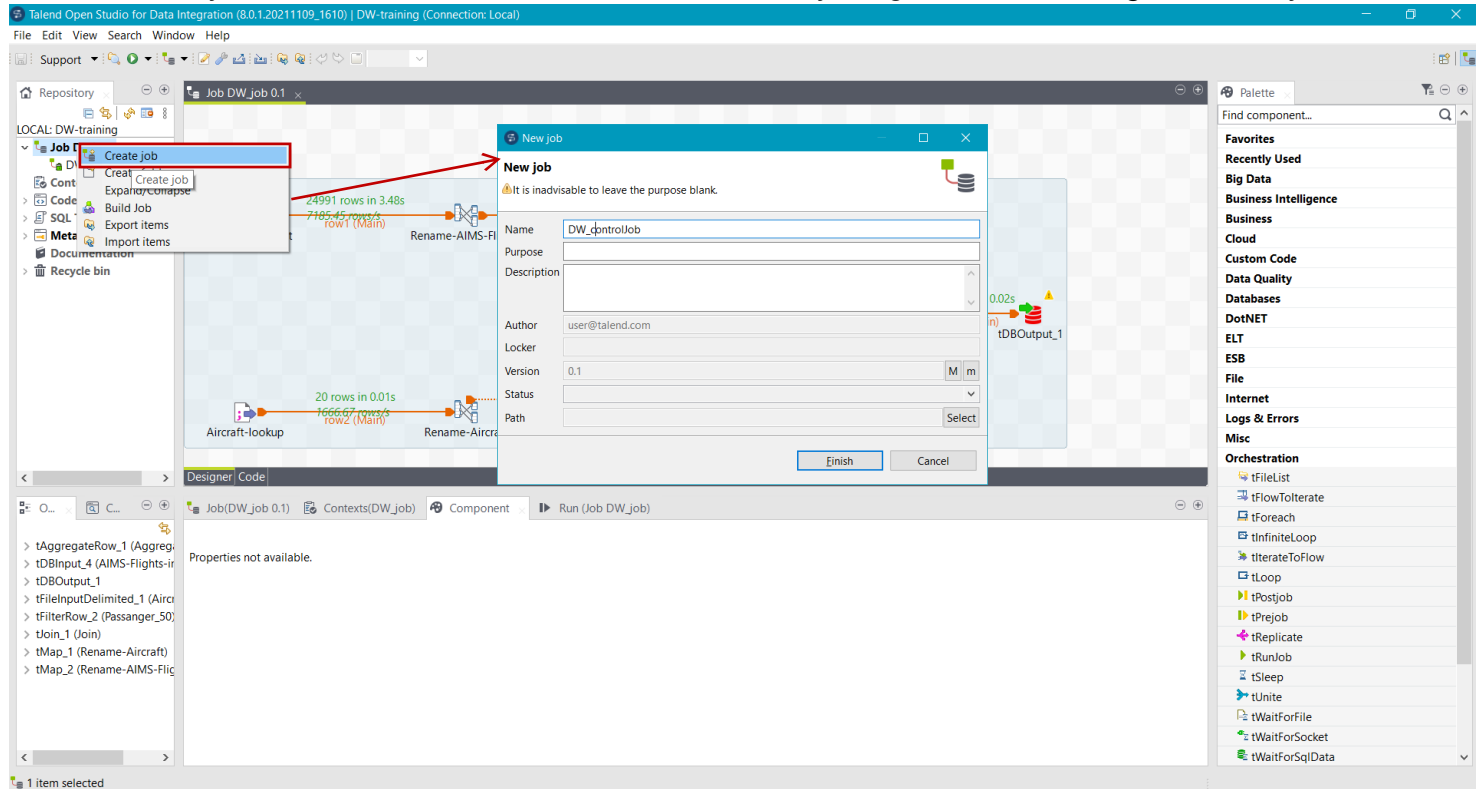
This is an excellent way of debugging Jobs when you are working with a relatively small sample of data. However, debugging Jobs that handle large volumes of data in this mode is not recommended because displaying every row visually in the Studio console can slow down the execution of the Job. In this case, you can try the Java Debug mode¹⁰.

¹⁰ <https://help.talend.com/r/en-US/8.0/debug-jobs/java-debug-mode>

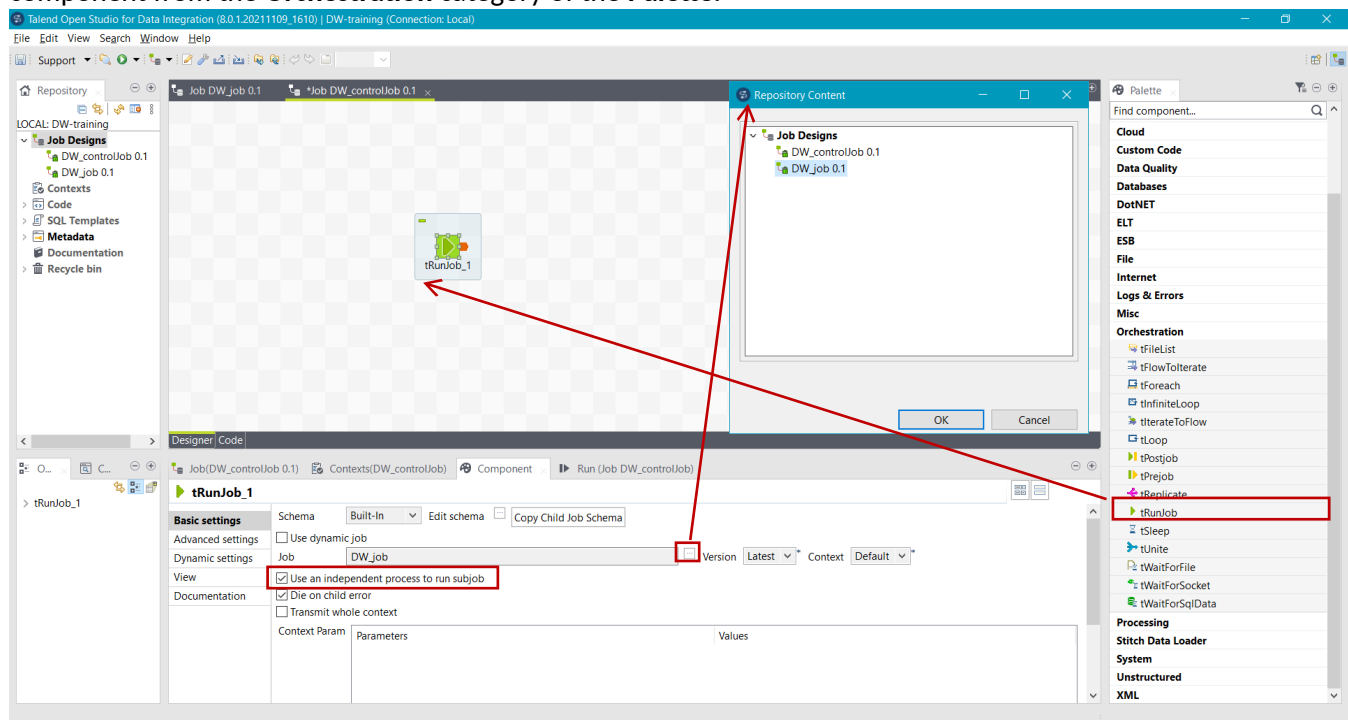
Part C: Create control flow (job) to execute an ETL process

Once we have our data flow(s) created, we need to manage their execution, i.e., to define the order in which they will be executed, in order to respect dependencies among them (e.g., first loading dimension tables then the fact table). To that end, as explained in Part A, we will create a simple control flow.

The control flow job is created in the same manner as the data flow job (right click to **Job Designs** -> Create job).



We next need to connect the control flow job to the previously created data flow job. To do this, we use the **RunJob** component from the **Orchestration** category of the **Palette**.



We first need to link the **RunJob** component to the data flow job by selecting it from the project repository (**DW_job**). In addition, we can also configure that the **RunJob** component uses an independent process to run the given data flow (**subjob**), by selecting this option in the component configuration. See the figure above.

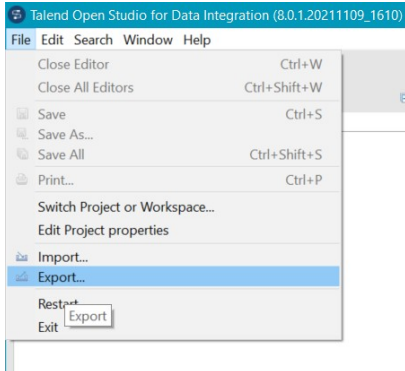
We run the control flow job in the same manner as the data flow job, i.e., by using the **Run** tab in the control panel. Similarly, we can add other data flows as subjobs using the **RunJob** component and connect them in order to orchestrate their execution, or use other orchestration components¹¹

¹¹ <https://help.talend.com/r/en-US/8.0/orchestration/orchestration> (Notice that not all the orchestrating components are available in the open source version of Talend).

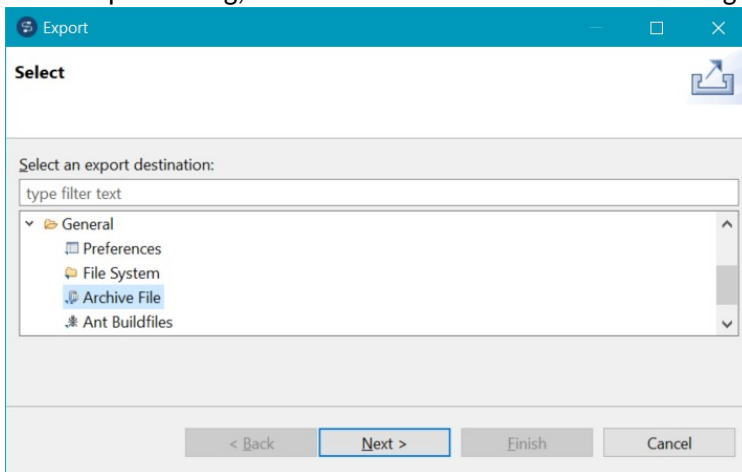
Appendix A - Exporting Talend project

To export your Talend project, in order to use it in another machine or for the final project delivery, you can follow these steps.

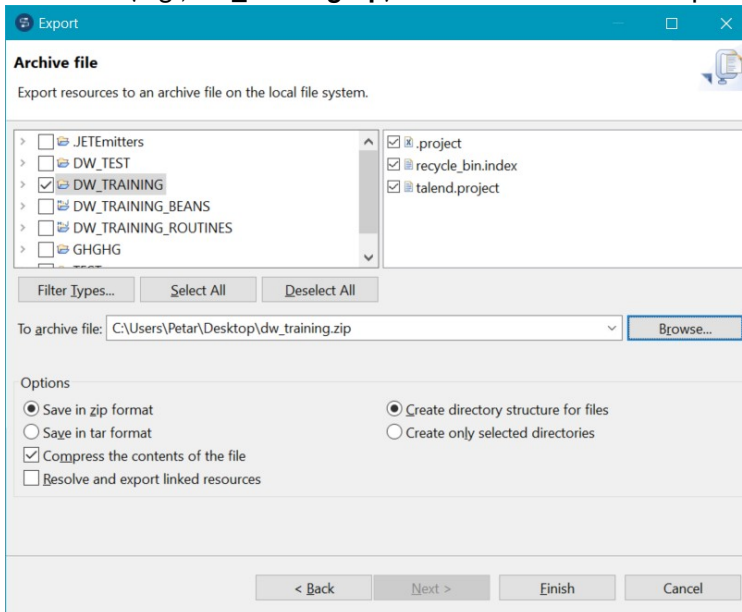
1. Under **File** menu click on **Export**



2. In the Export dialog, choose **Archive File** under **General** category, as a type of the export artifact. Then click Next.



3. Select the Talend project you want to export, choose the destination folder (**Browse...**) and give it a name with the extension (e.g., **dw_training.zip**). Be sure that the other options are as in the figure below. Then click Finish.



Your project is exported in a .zip archive file, and as such can be imported into another Talend installation¹²

¹² Import project: <https://help.talend.com/r/en-US/8.0/studio-user-guide-open-studio-for-data-integration/importing-local-projects>